Local network growth: How simple rules drive network complexity

Alexei Vazquez

July 27, 2026

To the ICTP, where it all started...

and Nodes & Links, for a renaissance.

Preface

In the year 2000 I was at the International Centre for Theoretical Physics (ICTP), located in Trieste, Italy. I had left my home country, Cuba, and gone there, with many dreams and much nostalgia. At the ICTP I indulged myself in the library, the scientific meetings, the computer cluster, and new friends. Then, on May 22–25, 2000, I attended the Research Workshop on Graph Theory and Statistical Physics. There was a presentation by an obscure Hungarian physicist by the name of Albert-László Barabási: "Emergence of scaling in complex networks." A small conference room and few attendees. Yet, it was eye-opening, a fascinating introduction to the properties of real networks and a hypothesis for their evolution. When I left that room I was charged, full of ideas. Two weeks later I had my first paper on network science: "Knowing a network by walking on it: emergence of scaling" (Vazquez, 2000). And right there, it all started.

Networks are everywhere, and so are the regularities they share. Whether one looks at the Internet, the World Wide Web, a society, a cell, or the scientific literature, the same statistical signatures recur: broad, often scale-free degree distributions; short distances between nodes; high local clustering that thins out in a regular hierarchy as degree grows; and systematic correlations between the degrees of neighboring nodes. That these features appear across systems with nothing else in common is one of the most striking findings of network science. Explaining them has been the field's central occupation for more than two decades.

The dominant explanation has rested on the preferential attachment mechanism, proposed by Réka Albert and Barabási (1999):

new nodes connect to existing ones in proportion to their degree, and a power law follows. It is a powerful idea, but taken literally it is a *global* rule, requiring each newcomer to weigh every node in the network against every other. Real web pages, proteins, and people do no such thing. They act on what is near them — the page a surfer reaches, the friend a friend introduces, the gene that gets copied. Preferential attachment is a macroscopic perception. In the same way, our perception of temperature is a macroscopic consequence of air molecules heating our skin. This book is about the microscopic mechanism behind preferential attachment and beyond: the regularities of real networks *emerge* from local rules. A growth rule that reaches just one step beyond a randomly chosen node, to one of its neighbors, already attaches preferentially by degree — and the same locality that produces preferential attachment produces the clustering hierarchy, the degree correlations, the community structure, and the multiplicity of shortest paths, all as a package. Coincidentally, I consolidated these ideas while working as a postdoctoral research associate in Barabási's group, then at the University of Notre Dame.

The chapters develop this thesis one mechanism at a time. After an introduction that fixes the measures and states the argument, I treat growth by search (citation and web graphs), by triadic closure (social networks), and by duplication (protein interaction networks) — the three local rules of my 2003 paper — and then carry the same principle into more recent work: the duplication–split structure of project schedules, the emergence of communities from local rules, and the boosting of shortest path multiplicity. Each chapter pairs a concrete local rule with a class of real networks it illuminates, and each ends where it began, with the conclusion that locality is the common cause.

The book is intended for anyone curious about how complex systems organize themselves — the interested general reader as much as the specialist. Readers drawn to the *general* principle will find it running through every chapter. Mathematicians, computer scientists, and physicists, in particular, will recognize a single idea applied again and again: that a handful of simple local rules, with no global coordination, is enough to account for the statistics that recur across otherwise unrelated networks; the closing chapters on the emergence of communities and on the redundancy of shortest

paths push that idea furthest. Readers drawn instead to a *particular* system will find it given a chapter of its own. Social scientists will recognize their networks of acquaintance and collaboration in the friends-of-friends mechanism of Chapter 3; biologists will find the evolution of protein interaction networks through gene duplication in Chapter 4; project managers and engineers will find the anatomy of project schedules — and the surprising irrelevance of the critical path — in the activity-specialization model of Chapter 5; and those who study the web, citation, or the spread of information online will find the search-and-connect rule of Chapter 2. Graduate students entering the field will find the models developed in full. Whatever the entry point, the core ideas are told in plain language, and the more technical calculations are set in boxes that can be skipped without losing the thread. I have tried to keep the analysis self-contained, favoring mean-field arguments that make the mechanisms transparent over exhaustive rigor, and to let the empirical data set the agenda throughout. Readers will find the degree denoted k, the elements of a graph called nodes, and their relations called links, following the now-standard usage of network science.

Much of this material grew out of collaborations and conversations over many years, with special mentions to Alessandro Vespignani, Yamir Moreno, Romualdo Pastor-Satorras, Amos Maritan, Alessandro Flammini, Albert-László Barabási, Sara Nadiv Soffer and Christos Ellinas. The original research papers on which the chapters draw are cited at the appropriate points and collected in the bibliography. Parts of the prose, and several of the figures, were prepared and edited with the help of Anthropic's Claude (Opus 4.8). I am grateful to the colleagues who shaped these ideas, and to the reader who, I hope, will come away persuaded that complex networks need not be designed — that they can grow themselves, one local decision at a time.

Alexei Vazquez

Contents

1 Introduction **1**
1.1 Networks everywhere 1
1.2 Some basic measures 3
1.3 From global to local rules 9
1.4 What the data demand 10
1.5 Plan of the book 11

2 Search and connect **13**
2.1 Finding what to link to 13
2.2 Surfing as a random walk 14
2.3 An effective "rich get richer" 16
2.4 The degree distribution 17
2.5 Cliquishness and correlations 18
2.6 Two ways to search 20
2.7 Summary 22

3 Friends of friends **23**
3.1 Triadic closure as a growth rule 23
3.2 Potential links 24
3.3 Degree distribution and clustering hierarchy 26
3.4 Connecting nearest neighbors 28
3.5 Why social networks are assortative 29
3.6 Summary 30

4 Gene duplication **31**
4.1 Networks that copy themselves 31
4.2 Duplication and divergence, in their simplest form . 32

4.3 Cliquishness from self-interaction 34
4.4 When duplication and divergence are coupled 35
4.5 Disassortative correlations 37
4.6 The unifying principle 38
4.7 Summary 39

5 Activity specialization **41**
5.1 Projects as growing networks 41
5.2 The duplication–split rule 42
5.3 Predecessors and successors: identical power laws . . 43
5.4 The vanishing critical path 44
5.5 Real project schedules 46
5.6 Delay risk as percolation 46
5.7 Summary 47

6 Emergence of communities **49**
6.1 Communities nobody designed 49
6.2 The Ramsey community number 50
6.3 Local rules grow communities 52
6.4 When globality fails: the Barabási–Albert model . . 55
6.5 How local must "local" be? 56
6.6 Summary 57

7 Boosting shortest path multiplicity **59**
7.1 The redundancy of shortest paths 59
7.2 Defining the multiplicity 60
7.3 The baseline: random networks 61
7.4 Local rules boost the multiplicity 62
7.5 Cycles and the boost 62
7.6 Two sides of one coin 64
7.7 Summary 65

8 Outlook **67**
8.1 One principle, many faces 67
8.2 What remains to be done 69

Chapter 1

Introduction

1.1 Networks everywhere

Look closely at almost any complex system and you will find, hiding inside it, a network. A network, in the sense used throughout this book, is nothing more exotic than a collection of things together with a record of which pairs of them are connected. The "things" are drawn as dots and called *nodes*. A connection between two of them is drawn as a line and called a *link*. That is the whole vocabulary needed to begin: a network is dots and lines, nodes and links.

What is remarkable is how many different systems fit this simple picture. The Internet is a network whose nodes are routers and whose links are the cables between them. The World Wide Web is a network of pages connected by the hyperlinks you click. A society is a network of people connected by acquaintance, friendship, collaboration, or contact. A living cell contains a network of proteins connected when two of them physically interact, and another of chemicals connected by the reactions that turn one into another. A project is decomposed into discrete activities and the logical dependencies between them. The scientific literature is a network of papers connected by citations, and of authors connected when they write a paper together (Newman, 2001, 2003; Albert and Barabási, 2002; Dorogovtsev and Mendes, 2002). These systems share almost nothing on the surface — one is silicon and glass, another is flesh

and acquaintance, another is ink and ideas — and yet, viewed as networks, they turn out to look astonishingly alike.

That likeness is the subject of this book. Over the past few decades, as researchers gathered maps of real networks ranging from a handful of nodes to billions, the same handful of features kept reappearing, no matter the system. Four of them organize everything that follows, so it is worth meeting them in plain language before we make them precise.

The first is the *small-world* property (Watts and Strogatz, 1998). In a real network you can usually get from any node to any other in only a few hops, even when the network is enormous — the famous "six degrees of separation" between any two people on Earth is one example. At the same time these networks are highly *cliquish*: your friends tend to be friends of one another, far more than chance would predict. Short hops everywhere, plus dense little neighborhoods: that combination is the small world.

The second is the *scale-free* property (Barabási and Albert, 1999; Faloutsos et al., 1999). If you count how many connections each node has, you find that most nodes have very few, while a small number — the *hubs* — have a huge number. Hollywood has a few actors who have worked with almost everyone and a multitude who have appeared in one film; the Web has a few pages everybody links to and billions that are barely linked at all. There is no "typical" number of connections in such a network, which is what the term *scale-free* is meant to capture.

The third, more subtle and the real connective tissue of this book, is that real networks carry systematic *correlations* and a *hierarchy* (Pastor-Satorras et al., 2001; Vazquez et al., 2002; Ravasz and Barabási, 2003; Vazquez, 2002). Whom a node connects to is not random with respect to how connected those neighbors are, and the cliquishness around a node depends in a regular way on how many connections the node has. We will make both of these ideas concrete in a moment.

The fourth, more recent and pitched at a larger scale than the rest, is that real networks divide into *communities*: groups of nodes densely connected among themselves but only loosely tied to the rest. People cluster into social circles, web pages into topics, proteins into functional modules — so that a network often looks less like

one seamless fabric than a patchwork of tight-knit groups stitched loosely together. Detecting these communities, and explaining why they arise, has grown into a research field of its own (Girvan and Newman, 2002; Fortunato, 2010; Palla et al., 2005).

The central question of the book is not *whether* these patterns exist — the evidence is by now overwhelming — but *where they come from*. The answer we will develop, over and over in different settings, is that they need not be designed in. They appear all by themselves in networks that grow according to *local rules*: rules in which each newly arriving node decides where to attach using only what is right in front of it — a node it has already reached, and that node's immediate neighbors. No node ever needs to survey the whole network. That a global pattern can emerge from purely local decisions is, in one sentence, the thesis of this book.

1.2 Some basic measures

To talk about networks precisely we need a few simple measures. None of them is harder than counting; we introduce them gently here and will lean on them in every chapter. Throughout, we picture a network with n nodes. When it helps to be formal we use the *adjacency matrix* a_{ij}, a bookkeeping device that is simply 1 if nodes i and j are connected and 0 if they are not — a giant table of yes/no answers to "are these two linked?"

Degree and degree distribution. The *degree* of a node is its number of links — how many neighbors it has. We write it k, so a node with $k = 3$ has three connections. (In formulas, the degree of node i is $k_i = \sum_j a_{ij}$, which just says: add up the yes-answers in row i of the table.) The *degree distribution* $p(k)$ is the recipe for how degrees are spread across the network: it is the fraction of nodes that have exactly k links, equivalently the chance that a node picked at random has degree k.

A network is called *scale-free* when this distribution follows a *power law*,

$$p(k) \sim k^{-\gamma}, \tag{1.1}$$

where the symbol $\sim$ means "behaves like" for large k, and γ (gamma)

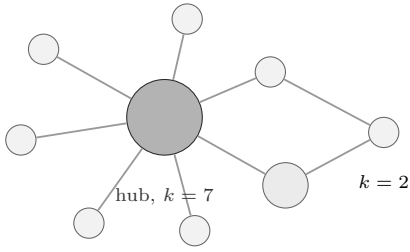


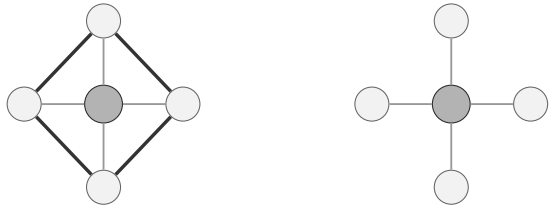
Figure 1.1: **Degree.** A node's degree is simply its number of links. Drawn here with node size growing with degree: a few hubs (such as the central one, $k = 7$) carry many links, while most nodes have only one or two. That imbalance — many small nodes, a few giant hubs — is the hallmark of a scale-free network.

is a number, the *exponent*, typically between 2 and 3 for real networks. A power law is just a rule that says "ten times bigger is some fixed number of times rarer." Unlike the bell-shaped curves of everyday statistics — human heights, say, which cluster tightly around an average — a power law has a long, heavy tail: extremely well-connected hubs are rare but nowhere near as rare as a bell curve would allow. This is why scale-free networks are hub-dominated, and why those hubs end up running much of the show (Figure 1.1).

Clustering coefficient. The *clustering coefficient* measures cliquishness: of all the pairs of your neighbors, what fraction are themselves connected? If you have four friends and four of the possible six friendships among them actually exist, your clustering coefficient is $4/6 \approx 0.67$. Formally, if node i has k_i neighbors and e_i of the possible links among those neighbors are present, then

$$c_i = \frac{2e_i}{k_i(k_i - 1)}, \tag{1.2}$$

since $k_i(k_i - 1)/2$ is the number of pairs that *could* be linked. A clustering coefficient near 1 means a tight clique; near 0 means your contacts are strangers to each other. Averaged over all nodes we write $\langle c \rangle$, where the angle brackets $\langle \cdot \rangle$ are standard shorthand for "average value." The striking empirical fact (Watts and Strogatz,



(a) high: $c \approx 0.67$ (b) low: $c = 0$

Figure 1.2: **Clustering coefficient.** The clustering of the central node (gray) is the fraction of its neighbor pairs that are themselves linked. (a) Four of the six neighbor pairs are linked (the bold links), so $c \approx 0.67$ — a cliquish neighborhood. (b) No neighbors are linked to each other, so $c = 0$.

1998) is that $\langle c \rangle$ in real networks is far larger than you would get by wiring the same nodes up at random — the quantitative fingerprint of those dense local neighborhoods (Figure 1.2).

Average distance. The *distance* between two nodes is the number of links on the shortest route between them — the fewest hops to get from one to the other. Averaged over all pairs of nodes, this distance grows very slowly (logarithmically, or even more slowly) as the network gets bigger. Slow growth of distance, together with high clustering, is precisely the small-world effect named above (Figure 1.3).

Degree correlations. Now to the more subtle pattern. Imagine following a link from a node out to one of its neighbors and asking: how connected is the neighbor I land on? If the network had no *degree correlations*, the answer would not depend on where I started. But real networks are correlated: in some, high-degree hubs tend to sit next to low-degree nodes; in others, hubs cluster together with other hubs. The cleanest way to see this is to record, for nodes of each degree k, the *average degree of their neighbors*, written $\langle k_{nn} \rangle(k)$ (the subscript "nn" stands for nearest-neighbors) (Pastor-Satorras

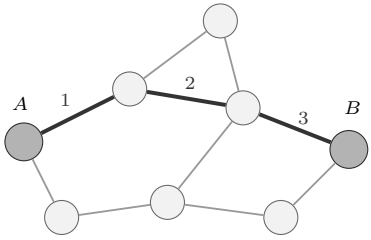


Figure 1.3: **Distance.** The distance between two nodes is the number of links on the shortest route between them. The shortest route from A to B here is three hops (in bold); the lower route is longer. In real networks these shortest distances stay small even as the network grows — the small-world effect.

et al., 2001):

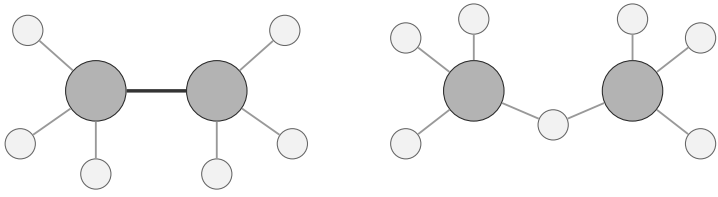
$$\langle k_{nn} \rangle(k) = \sum_{k'} k' \, p(k' \mid k). \tag{1.3}$$

Here $p(k' \mid k)$ is read "the probability that a link from a degree-k node lands on a degree-k' node," and the sum simply averages the landing degree k' weighted by how likely each is. Do not worry if the formula looks abstract; the next box pins down the one fact we need from it, and you can take that fact on faith and skip the box if you prefer.

It is worth knowing what this quantity would be if there were *no* correlations at all, because that is the yardstick against which real networks are judged.

The mathematics behind it: the no-correlation baseline for neighbor degree *(this box can be skipped)*

If degrees are uncorrelated, the only reason a link is more likely to land on a particular node is that the node has more links to land on. The chance of arriving at a node of degree k' is then proportional to k' times how common such nodes are, $k'p(k')$. Averaging the landing degree k' with this weight (and dividing by the total weight



(a) assortative (b) disassortative

Figure 1.4: **Degree correlations.** (a) In an *assortative* network, hubs tend to link to other hubs ("birds of a feather"). (b) In a *disassortative* one, hubs avoid each other and link instead to low-degree nodes ("opposites attract"). Social networks lean assortative; technological and biological ones, disassortative.

to normalize) gives

$$\langle k_{nn} \rangle_{\text{unc}} = \frac{\sum_{k'} k' \cdot k' p(k')}{\sum_{k'} k' p(k')} = \frac{\langle k^2 \rangle}{\langle k \rangle}, \tag{1.4}$$

where $\langle k \rangle$ is the average degree and $\langle k^2 \rangle$ the average of the squared degree.

The result, Eq. (1.4), is the key takeaway: in a network with no degree correlations, the average neighbor degree is the same flat number for every node, regardless of that node's own degree. So when we plot $\langle k_{nn} \rangle(k)$ against k for a real network and the curve is *not* flat, the tilt tells us the kind of correlation present. If the curve slopes *down* — well-connected nodes have poorly-connected neighbors — the network is called *disassortative*; this is typical of technological and biological networks. If it slopes *up* — the well-connected keep company with the well-connected — the network is *assortative*, as social networks tend to be (Newman, 2002). Assortative and disassortative are simply the "birds of a feather" and "opposites attract" of network connectivity (Figure 1.4).

Clustering hierarchy. Finally, we can combine the last two ideas by asking how cliquishness depends on degree. Sorting nodes by

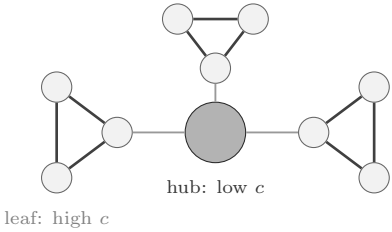


Figure 1.5: **Clustering hierarchy.** Low-degree nodes sit inside tight clusters where everyone is connected (high clustering, the dark triangles), while a hub bridges many such clusters. The hub's own neighbors — one drawn from each cluster — do not know each other, so the hub itself has low clustering. Cliquishness therefore falls as degree rises.

their degree k and averaging the clustering coefficient within each group gives $\langle c \rangle(k)$. In a wide range of real networks this quantity falls off as a power law,

$$\langle c \rangle(k) \sim k^{-1}, \tag{1.5}$$

meaning the more connections a node has, the *less* cliquish its neighborhood. This is the *clustering hierarchy*, and it has an intuitive reading (Ravasz et al., 2002; Ravasz and Barabási, 2003): ordinary low-degree nodes sit inside tight little communities where everyone knows everyone, while the big hubs act as bridges between many such communities, so their own neighbors — drawn from all over — mostly do not know each other. A small-town local has a dense circle of mutual friends; a globe-trotting connector knows thousands of people who have never met (Figure 1.5).

A word of caution ties the last two measures together, for they are not as independent as they look. The clustering hierarchy and the degree correlations are, in part, two views of the same thing. To see why, take a hub in a disassortative network: its neighbors are mostly low-degree nodes, and a low-degree node has too few links to be connected to many of the hub's other neighbors. The hub is therefore *forced* to have low clustering, not because of any designed hierarchy but simply because of whom it is attached to. Some

of the observed decay of $\langle c \rangle(k)$ with degree is thus an automatic by-product of the degree correlations, not independent evidence of modular organization. Indeed, when one strips out this bias and measures the clustering that remains, it is often roughly constant, or decays only very weakly, with degree (Soffer and Vazquez, 2005). The lesson is to read these measures together rather than in isolation — a theme that recurs throughout the book, where a single local rule will be seen to set the clustering hierarchy and the degree correlations *at once*, as facets of one mechanism.

1.3 From global to local rules

With the measures in hand we can state the puzzle this book exists to solve. The modern story of network growth opens with the model of Barabási and Albert (1999), usually abbreviated BA after its authors. It rests on two ideas. The first is *growth*: networks are not built all at once but accumulate nodes over time, the way the Web gains pages or the literature gains papers. The second is *preferential attachment*: a new node prefers to link to nodes that are already well connected, with a probability proportional to their degree — the "rich get richer" of networks. From these two ingredients a power-law degree distribution follows reliably, and preferential attachment has rightly become the textbook explanation for why so many networks are scale-free.

Two difficulties, though, drive everything that follows. The first is about *mechanism*. Preferential attachment, read literally, is a *global* rule: to attach in proportion to degree, a newcomer would have to know the degree of every node in the entire network and weigh them all against each other. That is implausible for the systems we care about. A scientist writing a new paper has read a few dozen others, not the millions in existence; a new gene is a copy of one existing gene, ignorant of the rest of the genome; a newcomer to a town meets a handful of people, not the whole population. Real nodes act *locally*. So if preferential attachment is not literally true, why do so many networks behave as if it were?

The second difficulty is that the BA model, for all its success with the degree distribution, does not by itself reproduce the other patterns we just met. Its clustering fades away as the network grows,

it shows none of the clustering hierarchy of Eq. (1.5), and its degree correlations are weak (Albert and Barabási, 2002; Vazquez, 2003). Something is missing.

The resolution pursued throughout this book is to throw out the global rule and replace it with local ones, then watch what *emerges*. By a *local rule* we mean a way of growing the network in which a newcomer reaches some node — chosen at random, or arrived at by surfing along links, or copied from an existing node — and then connects to that node and perhaps to a few of its immediate neighbors, using nothing but that local information. The recurring and rather beautiful discovery, established for several such rules (Vazquez, 2003), is that

> *an effective "rich get richer" behavior arises all on its own from growth by local rules.*

The reason, which we will meet many times, is almost a slogan: a well-connected node is, by definition, a node that many local explorations will bump into, because it has many links through which it can be reached. The preference for hubs need not be legislated; it is a free by-product of poking around the network locally. And because the very same local moves also create triangles (your friend introduces you to their friend) and link nodes of unlike degree, the same rules that produce the power law deliver, in one package, the clustering hierarchy of Eq. (1.5) and the degree correlations of Eq. (1.3). Three patterns that look independent turn out to share a single root.

1.4 What the data demand

Any explanation has to answer to measurements, so it is worth recording up front what those measurements actually look like across the systems in this book. When one plots the degree-resolved clustering $\langle c \rangle(k)$ and the average neighbor degree $\langle k_{nn} \rangle(k)$ for real networks, two broad families appear (Vazquez et al., 2002; Vazquez, 2003, 2002).

The first family — the maps of the Internet (both at the level of routers and at the level of the service providers, the "autonomous systems," that own them), the World Wide Web, the Gnutella file-sharing network, and the protein interaction network of a cell —

shows clustering that falls with degree, the hierarchy of Eq. (1.5), together with a neighbor degree that also falls with degree. These networks are *disassortative*: their hubs preside over peripheries of poorly-connected nodes rather than huddling together. This is the signature of technological and biological networks (Pastor-Satorras et al., 2001; Maslov and Sneppen, 2002; Jeong et al., 2001).

The second family — scientific collaboration and co-authorship networks — behaves oppositely on the correlation measure: the neighbor degree *rises* with degree, so these networks are *assortative*, with prolific, well-connected authors tending to write with one another (Newman, 2002). The contrast between the two families is instructive, because it shows that whatever mechanism we propose must be able to produce *either* sign of correlation, depending on the microscopic move — and, as the chapters will show, local rules can do exactly that.

These regularities are robust, they hold across many orders of magnitude in network size, and for years their shared origin was a mystery. Supplying that origin, one mechanism and one class of networks at a time, is the work of this book.

1.5 Plan of the book

The chapters proceed from mechanism to application, each built around one concrete local rule and a family of real networks it illuminates. The reader need not take any chapter on faith: in each, a simple local rule is stated, its consequences are worked out (with the heavier algebra confined to boxes that can be skipped), and the predictions are held up against real data.

- **Search and connect.** A newcomer learns of existing nodes by exploring the network — following references, surfing hyperlinks, performing a random walk — and connects to what it finds. This is the natural model for searchable, directed networks such as the citation network and the World Wide Web, and it is where we show in detail how the simple act of surfing manufactures an effective "rich get richer."

- **Friends of friends.** A newcomer connects to a node and then to some of that node's neighbors — you meet a person,

then their friends. This rule for social network growth creates triangles directly, and it generates the assortative correlations and the cliquishness hierarchy of social networks.

- **Gene duplication.** Biological networks grow by copying: a node is duplicated together with its connections, after which links are gained and lost. Duplication is an intrinsically local move, and it explains the disassortative, hierarchical structure of protein interaction networks.

- **Activity specialization.** The schedule of a large project is a network of activities grown by a local refinement rule: each activity is either duplicated into parallel copies or split into specialized sequential parts. The same copying mechanism gives power-law degrees, while making the celebrated "critical path" vanish as projects grow and relocating delay risk in the network as a whole.

- **Emergence of communities.** Local rules do not only set degrees and correlations; they also organize nodes into communities — tight groups with many links inside and few outside. We examine how this mesoscale structure arises from the same local dynamics, even when every node is otherwise identical.

- **Boosting shortest path multiplicity.** Finally we turn to the redundancy of routes through a network — how many equally short paths join two nodes — and show that local growth multiplies them, tying this redundancy to the very community structure of the previous chapter.

A single thread runs through all of them. In each case the apparently global features of a network — its scale-free degrees, its small-world distances, its clustering hierarchy, its degree correlations, its communities, its redundant paths — turn out to be the collective, emergent consequence of simple rules that no node applies with more than local knowledge. That is the sense in which the title of this book should be read: not networks built by a designer, but networks that grow themselves, one local decision at a time.

Chapter 2

Search and connect

2.1 Finding what to link to

The introduction left us with a puzzle. The "rich get richer" rule explains why so many networks are scale-free, but as a literal recipe it asks the impossible: a newcomer would have to know the degree of every node in the network in order to favor the well-connected ones. In this chapter we solve the puzzle for an important family of networks — the *searchable* ones — by replacing that impossible global rule with the most natural local act a newcomer can perform. It *searches* the network, moving from node to node along the links, and connects to what it happens to find.

The systems to keep in mind are collections of documents that point at one another. In the *citation network*, each node is a published article and each link is a reference from one article to another. In the World Wide Web, each node is a page and each link is a hyperlink. These are *directed* networks: a reference, like a hyperlink, points one way. An article cites an older one, not the reverse; a page links to another, which need not link back. Because direction matters, each node now has two kinds of degree. Its *in-degree*, written k^{in}, counts the links pointing *at* it — how many times a paper has been cited, how many pages link to a given page. Its *out-degree*, written k^{out}, counts the links pointing *out* of it — how many references a paper makes, how many hyperlinks a page contains. It is the

in-degree that interests us, because in real document networks it is the in-degree — the count of incoming attention — that is broadly distributed, with a few hugely-cited classics and a long tail of papers cited once or never.

How does a new document come to point at the old ones? Not by consulting a global popularity ranking. An author writing a new paper finds references by reading a few papers and following *their* reference lists, and then the references of those, and so on (Price, 1965). A web author discovers pages by surfing from link to link and by typing queries into a search engine (Brin and Page, 1998). In every case the information used is strictly local — a document one has reached, and the documents it points to. The claim of this chapter is that this local searching, repeated by every newcomer, adds up to exactly the "rich get richer" behavior that the introduction had to assume by fiat, and that it brings the power-law degrees, the clustering hierarchy, and the degree correlations along with it (Vazquez, 2001, 2003).

2.2 Surfing as a random walk

To turn "surfing" into something we can analyze, we model the surfer as a *random walker* on the network. A random walk is just what it sounds like: start somewhere, and step to a randomly chosen neighbor, then to a neighbor of *that* node, and so on. Our surfer does this with a twist. From whatever page it is on, with probability q_e it follows one of that page's outgoing hyperlinks, picked at random; with the remaining probability $1 - q_e$ it gets bored, stops following links, and jumps to a completely random page elsewhere. The single dial q_e, a number between 0 and 1, sets how patiently the surfer explores along links before restarting from scratch.

Now ask: in the long run, what fraction of its time does the walker spend at each page? Call that fraction the page's *visiting probability*, v_i for node i. A page is visited often if many links point to it and those linking pages are themselves visited often — a self-referential condition that turns out to have a clean form. The remarkable thing, worth flagging before any algebra, is that this visiting probability is *exactly the quantity that Google was built on*: it is the *PageRank* of the page, the measure of importance that ranks search results (Brin

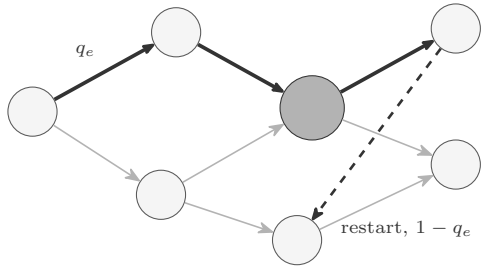


Figure 2.1: **Surfing as a random walk.** From the page it is on, the searcher follows one of the outgoing links with probability q_e (solid arrows) or, with probability $1 - q_e$, jumps to a page chosen at random (dashed arrow). A page reached along many routes — like the highlighted hub — is visited most often; this visiting frequency is the page's PageRank, and it grows in proportion to the number of links that point at the page.

and Page, 1998). So the local rule we are about to write is not a toy — it is the very process real search engines compute and then display to the authors who go on to create new links.

The picture to keep in mind is Figure 2.1. We would like to know how a page's visiting probability depends on its in-degree. Getting there exactly is hard, so we make the standard physicist's move of replacing fluctuating quantities by their averages — a so-called *mean-field approximation*, "mean" as in average. The next box carries this out; the one sentence to carry away is stated right after it, so the box can be skipped.

The mathematics behind it: from surfing to the visiting probability *(this box can be skipped)*

The visiting probability of node i obeys a balance condition: the walker arrives at i either by a random jump (probability $1 - q_e$, landing on any of the n nodes equally) or by following a link into i from some node j that points to it,

$$v_i = \frac{1 - q_e}{n} + q_e \sum_j a_{ij} \frac{v_j}{k_j^{out}}, \tag{2.1}$$

where $a_{ij} = 1$ when j points to i, and the factor $1/k_j^{out}$ accounts for j spreading its attention over its k_j^{out} outgoing links. This is the PageRank equation. To simplify it, replace the sum over the specific nodes pointing to i by an average. A node is reached along an incoming link from a node of out-degree k^{out} with probability $k^{out} p_{k^{out}} / \langle k^{out} \rangle$, and that link is then followed with probability $1/k^{out}$; the two factors of k^{out} cancel, so each incoming link contributes the same average amount $\Theta = \langle v \rangle / \langle k^{out} \rangle$. Since node i has k_i^{in} incoming links,

$$v_i = \frac{1 - q_e}{n} + q_e \, \Theta \, k_i^{in}. \tag{2.2}$$

The payoff is Eq. (2.2), and it is simple to state in words: a page's visiting probability grows in a straight line with its in-degree. Double the number of links pointing at a page and you roughly double how often the surfer lands on it. That straight-line relationship is the seed from which everything in this chapter grows.

2.3 An effective "rich get richer"

Being visited is not the same as being linked to — you read many pages you do not cite. So suppose that each time the surfer visits a page it creates a new incoming link to it only with some probability q_v. We can then ask how fast a page accumulates new links as a function of how many it already has. This rate is the network-growth analog of preferential attachment, and computing it is the crux of the chapter. The derivation is short; its conclusion is in the sentence after the box.

The mathematics behind it: the rate at which a page gains links *(this box can be skipped)*

Let ν_s surfers be active and ν_a new pages be added per unit time. The node count n and link count e grow as

$$\frac{\partial n}{\partial t} = \nu_a, \qquad \frac{\partial e}{\partial t} = \nu_s \, q_v \, \langle v \rangle \, n, \tag{2.3}$$

because each of the $\langle v \rangle n$ visits in a walk makes a link with probability q_v. Integrating fixes the average degree and hence $\Theta = \nu_a / (q_v \nu_s n)$.

The rate at which a page of in-degree k^{in} gains its next link is $A(k^{in}) = q_v\, v(k^{in})$; substituting Eq. (2.2) and the value of Θ,

$$A(k^{in}) = \frac{1}{n}\left[q_v(1-q_e) + q_e\,\frac{\nu_a}{\nu_s}\,k^{in}\right]. \tag{2.4}$$

Equation (2.4) is the central result, and it deserves to be read slowly. The rate at which a page attracts new links has two parts: a constant piece, the same for every page, which gives even a brand-new page a chance of being noticed; and a piece *proportional to its current in-degree*. That second piece is preferential attachment — the rich getting richer — but notice where it came from. The surfer never consulted a global ranking of pages; it only followed links, one at a time. The bias toward well-connected pages appeared on its own, for the simple reason that a page with many incoming links is a page the wandering surfer reaches by many routes. We have *derived* the rule the introduction had to assume. The same logic underlies the "redirection" trick of Krapivsky and Redner (2001), in which a newcomer points either to a random node or to whatever that node points to — a one-step walk that yields the same effective preference.

2.4 The degree distribution

A linear attachment rate is exactly the ingredient needed to produce a power law. Working out the resulting distribution of in-degrees requires solving a balance equation for how many pages sit at each in-degree, which we do in the box. Readers willing to take the outcome on trust can jump to the discussion after it.

The mathematics behind it: the in-degree distribution and its exponent *(this box can be skipped)*

Let $n_{k^{in}}$ be the number of nodes with in-degree k^{in}. It changes as nodes climb from one in-degree to the next by gaining links, with

new nodes injected at in-degree zero:

$$\frac{\partial n_{k^{in}}}{\partial t} = \nu_s \, A_{k^{in}-1} \, n_{k^{in}-1} - \nu_s \, A_{k^{in}} \, n_{k^{in}} + \nu_a \, \delta_{k^{in},0} \, , \tag{2.5}$$

the symbol $\delta_{k^{in},0}$ being 1 only at $k^{in} = 0$. For a network growing at a constant rate the ratio $\nu_s/\nu_a = \alpha$ is fixed and the shape of the distribution settles down, $n_{k^{in}}(t) = N p_{k^{in}}$. Solving for that stationary shape gives (with Γ the gamma function, a smooth generalization of the factorial)

$$p_{k^{in}} = \frac{1}{1+a} \, \frac{\Gamma[a(\gamma-1)+k^{in}]}{\Gamma[a(\gamma-1)]} \, \frac{\Gamma[(1+a)(\gamma-1)+1]}{\Gamma[(1+a)(\gamma-1)+k^{in}+1]}, \tag{2.6}$$

$$\gamma = 1 + \frac{1}{q_e}, \qquad a = \alpha \, q_v \, (1-q_e), \tag{2.7}$$

which for large in-degree is the power law

$$p_{k^{in}} \sim (k^{in})^{-\gamma}. \tag{2.8}$$

Two features of the result, Eq. (2.7), are worth dwelling on. First, the exponent γ is at least 2 — exactly the range seen in real citation and web networks. Second, and more surprising, γ depends *only* on q_e, the surfer's patience for following links. It does not depend on q_v, how often readers actually cite what they read, nor on the growth rate of the network. Whether authors are generous or stingy with citations changes how many links there are, but not the *shape* of the distribution; only the manner of *exploring* the network does that. As the surfer explores ever more deeply ($q_e \to 1$) the exponent drops to its minimum of 2, the most hub-dominated case; as the surfer restarts constantly ($q_e \to 0$) the exponent grows without bound and the heavy tail of hubs disappears.

2.5 Cliquishness and correlations

The same walk that sets the degrees also lays down triangles, so it predicts the clustering hierarchy of the introduction with no extra assumptions. The reason is simple to picture: whenever the surfer steps from a page to one of that page's neighbors and links to both, it has closed a triangle. A short calculation turns this picture into a

formula.

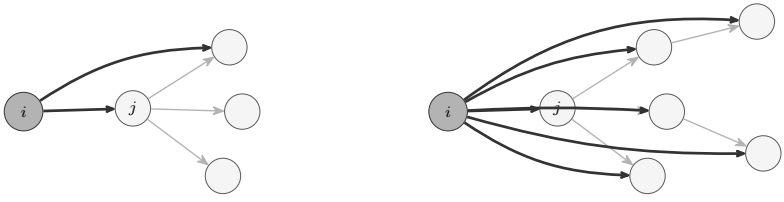
The mathematics behind it: how clustering falls with degree *(this box can be skipped)*

Counting the links e_i among the neighbors of node i created by the walk gives $\partial e_i/\partial t \approx (1+q_e)\,\partial k_i^{in}/\partial t$. Integrating this differential equation, assuming $e=0$ for $k=1$, we obtain $e=(1+q_e)(k-1)$. Inserting this into the definition of the clustering coefficient, $\langle c\rangle_k = 2e(k)/[k(k-1)]$, yields

$$\langle c\rangle_k \approx \frac{2(1+q_e)}{k}. \tag{2.9}$$

The message of Eq. (2.9) is that cliquishness is inversely proportional to degree: a node with twice the degree has roughly half the clustering. This is precisely the clustering hierarchy of Eq. (1.5), seen in the Internet, the Web, and the protein network — and here it falls out for free from growth by searching. The intuition is that a node gains a triangle every time the surfer steps from it to a neighbor, so its triangles grow in step with its degree, while the number of *possible* triangles grows much faster, like the square of the degree; their ratio, the clustering, therefore falls off as one over the degree.

The degree correlations are governed by the same single dial q_e. Numerical simulation of the model shows that when q_e is small the network looks essentially uncorrelated — the average neighbor degree is flat, as in the baseline of Eq. (1.4) — while as q_e grows the average neighbor degree develops the downward slope of a *disassortative* network, hubs sitting amid poorly-connected pages. This matches the real Web, whose measured exponent $\gamma \approx 2.1$ corresponds through Eq. (2.7) to a patient surfer with $q_e > 0.5$, squarely in the disassortative regime. One knob, q_e, thus fixes the degree-distribution exponent, the cliquishness hierarchy, and the sign of the correlations all at once — tying together what the introduction presented as three separate facts.



(a) random-walk model (b) recursive-search model

Figure 2.2: **Two ways to search, when a new node i joins.** (a) In the random-walk model, i links to a randomly chosen node j and then takes a single step to one of j's neighbors, linking to it as well. (b) In the recursive-search model, i follows *all* of j's links, and then their links in turn, linking to every node its widening search reaches. New links are drawn in bold, pre-existing links thin and gray.

2.6 Two ways to search

It helps to see two concrete versions of the searching rule, because they behave differently and bracket the range of possibilities. Figure 2.2 contrasts them.

The random-walk model. The simplest version uses a single surfer that always links what it visits and starts a fresh search from a new node each time it stops. Stated as a recipe, beginning from a couple of connected nodes, repeat:

- *Add.* Create a new node with one link pointing to a randomly chosen existing node.
- *Walk.* Having just linked to a node, then with probability q_e also link to one of that node's neighbors; otherwise go back to the add step.

Here the surfer takes a single step along the network before possibly stopping. This model produces the power-law tail in Eq. (2.8) and the clustering law of Eq. (2.9) (the latter in very good agreement with simulation), while the predicted exponent is a slight overestimate — the price of the mean-field approximation in the box above.

The recursive-search model. A more thorough searcher does not stop after one step but follows *all* the links of each node it reaches, then all the links of those, exploring an ever-widening territory (Vazquez, 2001). Specifically,

- *Add.* Create a new node with one link pointing to a randomly chosen existing node.
- *Walk.* If a link is created to a node in the network then with probability q_e a link is also created to each of its nearest neighbors. When no link is created go to the add rule.

This version has a richer behavior — in fact a genuine *phase transition*, a sudden change of character as the dial q_e is turned, of the kind familiar from water freezing into ice. The next box records the two cases that can be solved exactly and the location of the transition; the meaning is summarized afterward.

The mathematics behind it: the recursive searcher and its phase transition *(this box can be skipped)*

For $q_e = 0$ only the add step operates, every node is equally likely to be linked and $A(k^{in}) = 1/n$ independently of k^{in}. In this case the steady state solution $n_{k^{in}}(n) = n\, p_{k^{in}}$ of Eq. (2.5) is

$$p_{k^{in}} = 2^{-(k^{in}+1)}. \tag{2.10}$$

For $q_e = 1$ the recursion reaches every neighbor, the attachment rate becomes exactly linear, $A(k^{in}) = (1 + k^{in})/n$ — genuine preferential attachment, generated locally — and the steady state solution $n_{k^{in}}(n) = n\, p_{k^{in}}$ of Eq. (2.5) is

$$p_{k^{in}} = \frac{1}{(k^{in}+1)(k^{in}+2)} \sim (k^{in})^{-2}. \tag{2.11}$$

In between, we don't have an analytical solution. Yet, the average out-degree in the steady state is

$$\langle k^{out} \rangle = \frac{1 - \sqrt{1 - 4q_e}}{2q_e}, \tag{2.12}$$

which is finite only for $q_e \leq q_c = 1/4$; above this threshold it diverges.

The exact results bracket the behavior: a fast-decaying exponential when the searcher barely explores, and a clean power law with exponent 2 when it explores exhaustively. The new feature is the threshold q_c in Eq. (2.12). Below it the network has a well-defined average connectivity; above it the average degree runs away and the in-degree distribution locks onto the power-law exponent $\gamma = 2$, no matter the precise value of q_e. Simulations confirm this picture, with the threshold near $q_c \approx 0.5$.

The striking part is that scale-invariance — the power law — is not produced by carefully tuning q_e to one magic value, but holds across a whole range of q_e above the threshold. Physicists have a name for systems that sit automatically at such a critical, scale-free point without fine-tuning: *self-organized criticality*, the textbook example being a sandpile that, however you pour sand on it, keeps adjusting itself to the brink of avalanching. Here the slow addition of new pages and the fast exploration of the existing network play the roles of slow pouring and fast avalanching. Read this way, the ubiquity of scale-free citation and web networks is not a delicate coincidence but the generic, self-organized outcome of building a network by searching it.

2.7 Summary

Searchable networks grow by a rule with nothing global in it: a newcomer explores the existing network along its links and connects to what it finds. Modeling the explorer as a random walker — the very PageRank process that ranks web pages — we found that a page's visiting probability rises in a straight line with its in-degree, which converts the local act of walking into an effective "rich get richer." From that single mechanism follow a power-law in-degree distribution whose exponent $\gamma = 1 + 1/q_e$ is set purely by how the network is explored, a clustering hierarchy in which cliquishness falls as one over the degree, and tunable, generically disassortative degree correlations. Letting the searcher explore exhaustively turns the smooth picture into a self-organized critical transition with a robust exponent of 2. The puzzle of where preferential attachment "comes from," for this family of networks, is thereby dissolved: it comes from walking.

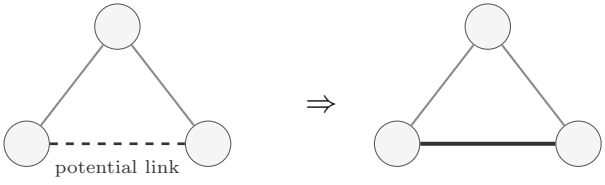
Chapter 3

Friends of friends

3.1 Triadic closure as a growth rule

The previous chapter built networks by exploration. This one builds them by introduction. Think about how you came to know the people you know. Only rarely did you befriend a complete stranger; far more often a friend introduced you to one of *their* friends. This everyday mechanism — a new tie forming between two people who already share a common acquaintance — has a name in sociology: *triadic closure*, because it "closes" a triangle by adding the third side to a path of two. Two people who both know a third are, in a sense, a triangle waiting to happen, and triadic closure is the event that completes it. Sociologists have long seen this as the engine of social structure (Granovetter, 1973). If most new links close triangles, then triangles pile up, and the network acquires the high cliquishness — the large clustering coefficient of the introduction — that sets social networks apart from random ones.

Figure 3.1 shows this elementary step. Notice that it is a thoroughly *local* rule. To introduce two friends you need to know only one node and its neighbors: who you know, and who they know. That is exactly the kind of local information this book is built on, and the question, as always, is what such a rule produces in the aggregate. The answer, worked out below following Vazquez (2003) and reviewed in Vazquez (2002), is that connecting friends-of-friends



(a) friend of a friend (b) triangle closed

Figure 3.1: **Triadic closure.** Two people (bottom) who already share a common friend (top) form a *potential link* — a friend-of-a-friend pair, drawn dashed. Closing the triangle turns that potential link into a real one, the elementary step by which the friends-of-friends rule builds social structure.

generates the whole suite of social-network signatures at once: an effective "rich get richer" and hence a power-law degree distribution, the clustering hierarchy in which cliquishness falls with degree, and — in pointed contrast to the technological and biological networks of the other chapters — *assortative* degree correlations, the tendency of well-connected people to know one another.

Models in this spirit have a history. Davidsen et al. (2002) introduced an acquaintance-network model (the DEB model, after its authors) in which people are repeatedly introduced to friends of friends, producing a small world with high clustering; Holme and Kim (2002) bolted an explicit triangle-forming step onto preferential attachment to give it tunable clustering. What follows organizes these ideas around a single bookkeeping device that makes the mechanism transparent — the *potential link*.

3.2 Potential links

The key idea is that, under triadic closure, a pair of nodes is not just "connected" or "not connected." A pair that already shares a common neighbor is a triangle waiting to close — a candidate in a way that an arbitrary pair of strangers is not. So we sort every pair of nodes into one of three states:

- *disconnected* (s): no link and no common neighbor — two strangers with no mutual acquaintance;
- connected by a *potential link* (p): no link yet, but at least one common neighbor — two friends-of-a-friend, poised to be introduced;
- connected by a *link* (e): already friends.

The *potential link* is the formal stand-in for "friend of a friend." We write k_i for the ordinary degree of node i (its number of friends) and k_i^* for its *potential degree*, the number of friends-of-friends it has — the number of potential links waiting at it.

Triadic closure is then a story about a pair changing state: a potential link (p) becomes a real link (e) when the triangle closes. To track this we need to know the rates at which pairs hop between the three states, and one relationship among those rates is the mechanical heart of the model. It says that potential links are born in proportion to a node's degree — introduce yourself to someone with many friends and you instantly acquire many friends-of-friends. The next box sets up the bookkeeping and draws out its main consequence; the conclusion is restated plainly afterward, so the box may be skipped.

The mathematics behind it: from the three-state bookkeeping to "rich get richer" *(this box can be skipped)*

Let $\nu_{x\to y}$ be the rate (per unit of network size n) at which a pair hops from state x to state y. Ignoring fluctuations, the average degree and potential degree of node i evolve as

$$\begin{aligned} n\,\frac{\partial k_i}{\partial n} &= \nu_{s\to e}\,\hat{k}_i + \nu_{p\to e}\,k_i^* - (\nu_{e\to s} + \nu_{e\to p})\,k_i, \\ n\,\frac{\partial k_i^*}{\partial n} &= \nu_{s\to p}\,\hat{k}_i + \nu_{e\to p}\,k_i - (\nu_{p\to s} + \nu_{p\to e})\,k_i^*, \end{aligned} \tag{3.1}$$

where $\hat{k}_i = n - k_i - k_i^*$ is the number of nodes that are neither friends nor friends-of-friends of i. The link between the two kinds of connection comes from the geometry of triangles: when a new node i attaches to an existing node j, it instantly shares the common neighbor j with *all* of j's neighbors, gaining a potential link to each.

Hence potential links are spawned in proportion to degree,

$$\nu_{s\to p} = \nu_{s\to e}\, k_i, \qquad \nu_{p\to s} = \nu_{e\to s}\, k_i. \tag{3.2}$$

Now make two modeling choices. First, neglect link deletion, $\nu_{e\to s} = 0$ — accurate for, say, a co-authorship network, where having written a paper together is a fact that cannot later become false. Second, make closing a triangle far likelier than linking two strangers,

$$\nu_{s\to e} = \frac{\mu_0}{n^2}, \qquad \nu_{p\to e} = \frac{\mu_1}{n}, \tag{3.3}$$

the stranger-linking rate carrying an extra factor of $1/n$. Substituting these into Eq. (3.1) collapses the system to

$$n\,\frac{\partial k_i}{\partial n} = \mu_0 + \mu_1\, k_i^*, \qquad n\,\frac{\partial k_i^*}{\partial n} = \mu_0\, k_i - \mu_1\, k_i^*. \tag{3.4}$$

The upshot of Eq. (3.4) is that a node's degree grows at a rate proportional to how many friends-of-friends it has — and a node with many friends has, through them, many friends-of-friends. So once again the rich get richer, but this time the preference is manufactured by triadic closure rather than by surfing: a well-connected person is simply involved in more potential introductions. As in the previous chapter, no node ever consulted a global ranking; the bias toward hubs is a free by-product of a purely local act.

3.3 Degree distribution and clustering hierarchy

Because the growth rate is linear in degree, the same reasoning that gave a power law in the last chapter applies here, and it simultaneously fixes the cliquishness. The box carries out both; the two results and their meaning are in the paragraph after it.

The mathematics behind it: the power law and the clustering hierarchy *(this box can be skipped)*

Integrating the linear system (3.4) gives power-law growth of both degrees with the age of the node,

$$k_i(n) = k_0 \left(\frac{n}{n_i}\right)^{\beta}, \qquad k_i^*(n) = k_0^* \left(\frac{n}{n_i}\right)^{\beta}, \tag{3.5}$$

where n_i is the network size when node i arrived and

$$\beta = \frac{\mu_1}{2}\left(-1 + \sqrt{1 + 4\,\frac{\mu_0}{\mu_1}}\right). \tag{3.6}$$

With nodes added at a constant rate, $P(n_i{=}n) = 1/n$, a node has degree above k only when it is old enough for Eq. (3.5) to exceed k; averaging over its birth time gives the cumulative distribution

$$P(k_i > k) = P\left[k_0 \left(\frac{n}{n_i}\right)^{\beta} > k\right] = \int_0^n \frac{dn_i}{n}\,\Theta\left[k_0 \left(\frac{n}{n_i}\right)^{\beta} - k\right] \tag{3.7}$$

with $\Theta[\cdot]$ the step function, 1 when its bracket holds and 0 otherwise. Differentiating it leaves a power-law degree distribution,

$$p_k = \frac{\partial P(k_i > k)}{\partial k} \sim k^{-\gamma}, \qquad \gamma = 1 + \frac{1}{\beta}. \tag{3.8}$$

For the clustering, the dominant way node i gains a triangle is a potential link closing: when the potential link from i to j (sharing neighbor k) becomes real, i gains both the neighbor j and a link j–k among its neighbors. Thus $\partial e_i/\partial n = \nu_{p\to e}\, k_i^* = \mu_1 k_i^*/n$, and integrating with Eq. (3.5),

$$\langle c \rangle_k = \frac{2e(k)}{k(k-1)} \approx \frac{2\mu_1}{k}. \tag{3.9}$$

Two familiar results emerge. First, Eq. (3.8): the degree distribution is again a power law. It is worth knowing exactly where the power law comes from — it comes from the proportionality in Eq. (3.2), the fact that potential links are spawned in proportion to degree. Were they instead spawned at a flat rate, the mechanism would still create triangles but the degree distribution would

decay exponentially, with no fat tail of hubs. Scale-freeness is the fingerprint of that degree-proportional spawning. Second, Eq. (3.9): cliquishness again falls as one over the degree — the clustering hierarchy of the introduction, now arriving by a third route (after search, and before duplication in the next chapter). Here its origin is especially vivid: every triangle a node owns is just a friend-of-a-friend pair that has been introduced.

3.4 Connecting nearest neighbors

The continuum theory is borne out by a minimal simulation, a growing version of the DEB rule (Davidsen et al., 2002). Starting from a single node and no links, repeat:

- with probability $1-u$, *add* a new node and link it to a randomly chosen node j — which by definition creates a potential link between the newcomer and every neighbor of j;
- with probability u, *close* a triangle: pick a potential link at random and turn it into a real link.

The single dial u is the rate of triangle-closing relative to adding newcomers. Feeding the mapping $\mu_0 = 1$, $\mu_1 = u/(1-u)$ into the formulas above gives the exponent in closed form, as the short box records.

The mathematics behind it: the exponent as a function of the closure rate *(this box can be skipped)*

With $\mu_0 = 1$ and $\mu_1 = u/(1-u)$, Eqs. (3.6) and (3.8) give

$$\gamma(u) = 1 + \frac{2(1-u)}{u}\left(-1 + \sqrt{1 + 4\,\frac{1-u}{u}}\right)^{-1}, \tag{3.10}$$

with limiting values

$$\gamma(0) = \infty, \qquad \gamma(1) = 2. \tag{3.11}$$

In words, the exponent is tunable by how eagerly the network closes triangles. When closure is frequent and newcomers rare

($u \to 1$) the exponent sinks to its minimum of 2, the most hub-dominated case; when newcomers arrive so fast that triangles never get a chance to close ($u \to 0$) the exponent blows up and the heavy tail of hubs is lost. Simulations up to a million nodes confirm the power-law degree distribution (the mean-field formula slightly overestimating the exponent, as in the previous chapter), and they confirm the clustering hierarchy, with cliquishness falling as a power of degree.

3.5 Why social networks are assortative

The most telling prediction is about the degree correlations — and here the friends-of-friends rule parts company with every other rule in this book. Measuring the average neighbor degree $\langle k_{nn} \rangle(k)$ in the simulated network, one finds it *rises* with degree: the well-connected are connected to the well-connected. The network is *assortative* (Newman, 2002) — "birds of a feather" — exactly as measured in real collaboration and co-authorship networks, and exactly opposite to the disassortative Internet and protein networks of the introduction.

Why the opposite sign? The reason becomes obvious once seen. Every new link here closes a triangle, joining two nodes that share a common neighbor. But a node that serves as the common neighbor of many pairs is, by definition, a high-degree node, and the two nodes it brings together are drawn from its own neighborhood, which is better connected than the network at large. So closure tends to wire together nodes of *similar*, and similarly high, degree. The very mechanism that the disassortative models of the other chapters lack — linking within a neighborhood rather than reaching out to a random or peripheral node — is what tips the correlations positive. This is the deeper lesson of the chapter, and of the review that develops it (Vazquez, 2002): the *sign* of the degree correlations is not an arbitrary fitting knob but a direct readout of the microscopic rule. Search and duplication reach outward, to the periphery, and so make hubs that shun one another; triadic closure works inward, within neighborhoods, and so makes hubs that keep each other's company.

3.6 Summary

Modeling social growth by triadic closure — a newcomer attaches to a node and then, over time, to that node's friends — and formalizing "friend of a friend" as a *potential link*, we obtain a tractable theory in which potential links are created in proportion to degree. From that one ingredient follow an effective "rich get richer," a power-law degree distribution with a tunable exponent $\gamma = 1 + 1/\beta$, a clustering hierarchy in which cliquishness falls as one over the degree, and — distinctively — assortative degree correlations. Connecting friends-of-friends thus reproduces, from purely local moves, the characteristic shape of social networks, and explains why the sign of their correlations differs from that of the technological and biological networks treated elsewhere in this book.

Chapter 4

Gene duplication

4.1 Networks that copy themselves

The two rules of the preceding chapters each had a newcomer reach into the existing network and *choose* where to attach — by surfing, or by closing a triangle. Biology offers a third rule, and in some ways a more primitive one: the network grows by *copying* a node that is already there. The newcomer does not pick its neighbors at all. It inherits them, wholesale, from a parent.

The setting is the inner machinery of the living cell. Genomes evolve in good part through *duplication* — a gene, or sometimes an entire genome, is accidentally copied — followed by *divergence*, the slow accumulation of mutations that make the two copies drift apart. That duplication is a major engine of evolution was argued classically by Ohno (1970). To turn this into a network statement we use the *protein interaction network*: each node is a protein (the product of a gene), and a link joins two proteins that physically interact — that bind to each other to do their job. Now watch what a gene duplication does to this network. Right after the gene is copied, the cell makes two identical proteins, so the new protein interacts with *exactly the same partners* as the old one: the duplicate node arrives already wired to all the neighbors of its ancestor (Figure 4.1). Mutations then add a few interactions and, more often, remove some, letting the twins diverge (Wagner, 2001; Solé et al., 2002). Both



(a) duplication (b) divergence

Figure 4.1: **Duplication and divergence.** (a) A gene duplication adds a copy i' (gray) wired to every interaction partner of the ancestor i, plus a link between the twins when the protein is self-interacting. (b) Divergence then prunes the now-redundant interactions (dotted), each partner kept by only one twin, so the copies drift apart.

steps are thoroughly local — the copy needs to know only its parent and its parent's neighbors, never the rest of the genome.

The same picture reaches beyond biology. A new web page is often made by copying another page's hyperlinks and then editing the list — duplication as copying, divergence as editing. Following Vazquez (2003) and Vazquez et al. (2003), this chapter shows that copy-and-diverge, like the rules before it, produces an effective "rich get richer," a power-law degree distribution, the clustering hierarchy, and — as actually measured in protein networks (Maslov and Sneppen, 2002) — *disassortative* correlations. It also reveals something new: a *multifractal* degree distribution, which we will unpack when we reach it, a fingerprint left by the act of inheritance.

4.2 Duplication and divergence, in their simplest form

Start with the simplest version, treating duplication and divergence as independent random events. The single idea that drives everything is this: a node gains a new link whenever one of its *neighbors* is duplicated, because the neighbor's copy inherits a link back to it. A well-connected node has many neighbors, so it has many chances

to be on the receiving end of a duplication — and there, already, is the rich getting richer. The box makes this precise and reads off the degree distribution; the conclusion follows it.

The mathematics behind it: degree growth and the power law from duplication *(this box can be skipped)*

In a mean-field (average-quantities) description, with increasing the number of nodes n the degree of node i evolves as

$$\frac{\partial k_i}{\partial n} = \nu_D\, k_i + \nu_C\,(n - k_i) - \nu_L\, k_i, \tag{4.1}$$

with ν_D, ν_C, ν_L the per-node rates of duplication, random link creation, and link loss. The first term is the key one: it is proportional to k_i because each of i's k_i neighbors, if duplicated, hands i a new link. Each duplication adds one node, so $\nu_D = 1/n$; writing $\nu_C = \mu_0/n$ and $\nu_L = \mu_1/n$ reduces the equation to

$$n\,\frac{\partial k_i}{\partial n} = \mu_0 + (1 - \mu_1)\, k_i, \tag{4.2}$$

whose growth rate is linear in the degree. Integrating it gives the degree of node i as the network grows,

$$k_i(n) = \left(k_i(n_i) + \frac{\mu_0}{1-\mu_1}\right)\left(\frac{n}{n_i}\right)^{\beta} - \frac{\mu_0}{1-\mu_1}, \tag{4.3}$$

where n_i and $k_i(n_i)$ are the network size and the node's degree at the moment it was added, and the exponent is

$$\beta = 1 - \mu_1 \qquad (\mu_1 < 1), \tag{4.4}$$

the condition $\mu_1 < 1$ keeping the network from emptying out. To turn this growth law into a degree distribution, note that node i has degree above k exactly when its age and birth degree make Eq. (4.3) exceed k,

$$P(k_i > k) = \Theta\left[\left(k_i(n_i) + \frac{\mu_0}{1-\mu_1}\right)\left(\frac{n}{n_i}\right)^{\beta} - \frac{\mu_0}{1-\mu_1} > k\right], \tag{4.5}$$

with $\Theta[\cdot]$ the step function, 1 when its bracket holds and 0 otherwise. Both its birth size n_i and its birth degree $k_i(n_i)$ are random: for nodes added at a constant rate n_i is uniform, $P(n_i{=}n) = 1/n$, while

a node inherits its ancestor's degree, so $P[k_i(n_i){=}k'] = p_{k'}$ in the stationary state. Averaging over both,

$$P(k_i > k) = \sum_{k'} p_{k'} \int_1^n \frac{dn_i}{n}\, \Theta\left[\left(k' + \frac{\mu_0}{1-\mu_1}\right)\left(\frac{n}{n_i}\right)^{\beta} - \frac{\mu_0}{1-\mu_1} > k\right] \tag{4.6}$$

Differentiating this cumulative distribution and taking $n \gg 1$ leaves, for nodes added at a constant rate, a power-law degree distribution

$$p_k = \frac{\partial P(k_i > k)}{\partial k} \sim \left(\frac{\mu_0}{1-\mu_1} + k\right)^{-\gamma}, \qquad \gamma = 1 + \frac{1}{1-\mu_1}. \tag{4.7}$$

The result is once again a power law, Eq. (4.7), and its exponent is controlled entirely by μ_1, the rate at which divergence removes links: the more aggressively the duplicates shed interactions, the steeper the tail and the fewer the giant hubs. As in the previous two chapters, the power law can be traced to a single term that grows in proportion to degree — here the duplication term of Eq. (4.2) — and without it the distribution would decay as a featureless exponential. Copying, all by itself, is a rich-get-richer engine.

4.3 Cliquishness from self-interaction

If divergence added and removed links purely at random, this model would make no triangles, and its cliquishness would fade to nothing in a large network. But biology supplies one specific local source of triangles. Some proteins are *self-interacting* — they bind to copies of themselves. When the gene for such a protein is duplicated, the ancestor and its new twin can bind to *each other*, because each still carries the self-binding site (Wagner, 2001). That single interaction, laid down at the moment of duplication, is the seed of a triangle. The short box turns this into the clustering law.

The mathematics behind it: the clustering hierarchy from self-interaction *(this box can be skipped)*

Let a duplicated self-interacting protein bond to its ancestor with probability q_v. Then whenever a neighbor of node i is duplicated, i gains not only the copy as a new neighbor but, with probability q_v, a link between two of its neighbors (the ancestor and the copy). Hence the number of links e_i between the neighbors of node i change as

$$\frac{\partial e_i}{\partial n} \approx q_v \frac{\partial k_i}{\partial n}. \tag{4.8}$$

Integrating, assuming $e_i = 0$ when $k_i = 1$, we get $e_i = q_v(k_i - 1)$. Inserting this result into the clustering definition,

$$\langle c \rangle_k = \frac{2e(k)}{k(k-1)} \approx \frac{2q_v}{k}. \tag{4.9}$$

So cliquishness again falls as one over the degree, Eq. (4.9) — the clustering hierarchy of the introduction, reached now by a third independent route. There is one telling refinement: because removing an interaction also destroys triangles, and the biggest hubs have the most triangles to lose, divergence makes the clustering of high-degree nodes decay a little *faster* than one over the degree. That slightly-too-steep falloff is a small signature that distinguishes duplication from the search and closure mechanisms.

4.4 When duplication and divergence are coupled

In a real protein network the two processes cannot truly be independent. The network does a job — it keeps the cell alive — so losing the wrong interaction can be lethal, and which links a duplicate keeps is under natural selection. The classical view of Ohno (1970) lets one redundant copy decay freely; the more modern *subfunctionalization* picture of Force et al. (1999) has the two copies divide the ancestor's duties between them, each keeping part, together keeping the whole. Either way the fates of the two copies' links are *coupled*: a shared interaction is far more likely to be retained by one copy

than dropped by both. To capture this, Vazquez et al. (2003) use the following rule. Starting from two connected nodes, repeat:

- *Duplication.* Pick a node i at random and add a copy i' joined to every neighbor of i; with probability q_v also add the link i–i' (the self-interaction case).
- *Divergence.* For each node j now joined to both i and i', delete one of the two redundant links (i, j) or (i', j), at random, with probability $1 - q_e$.

The divergence step touches only the *redundant* pairs just created, and that is what couples it to duplication. Genomic data showing that duplicates diverge soon after they are born (Wagner, 2001) justify treating divergence as finishing before the next duplication, so the node count n doubles as a clock. A related model with lopsided divergence is studied in Vazquez (2003).

Coupling makes the large-scale behavior considerably richer, in two ways worth previewing before the box. First, a single number — the retention parameter q_e — decides whether the network stays sparse or grows ever denser, with a sharp switch between the two regimes. Second, the degree distribution is no longer a plain power law but a *multifractal* one. A quick word on that term: a plain power law is described by one exponent, the same at every scale; a multifractal is described by a whole continuum of exponents, because different "moments" of the distribution (the average degree, the average of the squared degree, and so on) scale with different powers of the network size. The box records the transition and the moment exponents; the interpretation follows.

The mathematics behind it: the density transition and the multifractal spectrum *(this box can be skipped)*

Tracking the mean degree through one duplication event gives

$$\langle k\rangle(n+1) = \frac{n\langle k\rangle(n) + 2q_v + (2q_e - 1)\langle k\rangle(n)}{n+1} \tag{4.10}$$

which in the large-network limit switches behavior at $q_e = 1/2$: for $q_e < 1/2$ the mean degree saturates at the finite value $\langle k\rangle = 2q_v/(1 - 2q_e)$, while for $q_e > 1/2$ it grows without bound as $N^{2q_e - 1}$. A rate-equation treatment in the manner of Kim et al. (2002) gives

an attachment rate $A(k) = (q_v + q_e k)/n$ — linear in degree again — and predicts that the l-th moment of the degree distribution scales as $M_l \sim n^{\sigma_l(q_e)}$ with

$$\sigma_l(q_e) = l\, q_e + 2\left[\left(\frac{1+q_e}{2}\right)^l - 1\right]. \tag{4.11}$$

For each l there is a value q_l at which σ_l changes sign, with $q_1 = 1/2$ and $q_2 = 2\sqrt{3} - 3 \approx 0.46$.

The two messages of the box are these. First, Eq. (4.10) shows a genuine switch at $q_e = 1/2$: keep too few links and the network settles into a sparse steady state; keep enough and it densifies forever. Second, and more subtle, the moment exponents σ_l in Eq. (4.11) depend on l in a curved, *nonlinear* way — the hallmark of a multifractal. Where the earlier chapters gave clean single power laws, duplication gives a spectrum of exponents, and the reason is exactly the inheritance at its core: each newcomer copies the degree environment of its parent, injecting a kind of heritable heterogeneity that a single exponent cannot capture. The deletion of links from the copies is what converts that inheritance into multifractality.

4.5 Disassortative correlations

Simulations spanning network sizes from a thousand to a million nodes confirm the analysis (Vazquez et al., 2003; Vazquez, 2003): the moment exponents match the multifractal prediction; the clustering falls with degree as a power law with exponent at or above one (steepening as link loss grows, just as Eq. (4.9) and its caveat foretold); and the average neighbor degree *decreases* with degree, $\langle k_{nn}\rangle(k) \sim k^{-0.1}$. The duplication network is therefore *disassortative* — its hubs are surrounded by poorly-connected partners rather than by each other.

This is the correct sign. Disassortative mixing is precisely what Maslov and Sneppen (2002) measured in the protein interaction network of yeast, and it files duplication alongside the technological networks of Chapter 2 and against the social networks of Chapter 3. With this third mechanism in hand, the three rules of the book so

far bracket the empirical landscape, as the table summarizes (here β is the rate at which cliquishness falls with degree, and α the slope of the neighbor-degree curve — positive for assortative, negative for disassortative):

Mechanism	$\langle c \rangle_k \sim k^{-\beta}$	$\langle k_{nn} \rangle_k \sim k^{\alpha}$
Connecting neighbors	$0 < \beta < 1$	$\alpha > 0$
Random walk	$\beta = 1$	$\alpha \leq 0$
Duplication–divergence	$\beta \geq 1$	$\alpha < 0$

Every local rule produces a clustering hierarchy; what differs is the *sign* of the correlations, which sorts the rules into the assortative (social, closure-driven) and disassortative (technological and biological, reach-driven) classes seen in real data.

4.6 The unifying principle

Three different stories — surfing a network, closing a triangle, copying a node — have each delivered the same effective "rich get richer." That is no accident, and the duplication picture makes the reason simple enough to state in a single sentence.

Picking a node *at random* plays no favorites: every node is equally likely, hub or not. But none of the local rules in this book act on a randomly chosen node. They act on a *neighbor* of one — the page a surfer steps to, the friend a friend introduces, the partner a duplicated gene inherits. And the chance that a particular node i is the neighbor of a randomly chosen node is

$$\frac{k_i}{\sum_j k_j}, \tag{4.12}$$

which is exactly the linear preferential-attachment probability of the Barabási–Albert model (Barabási and Albert, 1999). High-degree nodes are over-represented among the neighbors of a random node for the most elementary of reasons: they belong to more neighborhoods, so more random nodes have them as a neighbor. Preferential attachment never has to be imposed. It is induced automatically the instant a growth rule takes one step out from a random node to one of its neighbors. The clustering hierarchy has

the same local root — a rule that makes one link to a node tends to make another among that node's neighbors — and the sign of the degree correlations depends on *which* neighbor gets chosen. Locality is the common cause, and the three mechanisms are three faces of it.

4.7 Summary

Biological networks grow by duplication: a node is copied along with its connections, and the copies then diverge by gaining and losing links. Because a node gains a link whenever one of its neighbors is duplicated, the rule builds in an effective "rich get richer" and hence a power-law degree distribution, while duplications of self-interacting proteins generate a clustering hierarchy that link loss steepens. Coupling divergence to duplication adds a density transition at $q_e = 1/2$ and a multifractal degree distribution born of inheritance, and the model reproduces the disassortative correlations measured in real protein networks. Together with the two previous chapters, duplication completes a simple and general principle: a growth rule that attaches to a *neighbor* of a randomly chosen node attaches preferentially by degree — and so preferential attachment, the clustering hierarchy, and degree correlations are all the generic harvest of locality.

Chapter 5

Activity specialization

5.1 Projects as growing networks

The duplication rule of the previous chapter was born in biology, but copying is not a biological privilege. This chapter applies the same idea to a very different system — the schedule of a large engineering project — and gets, in return, a surprising practical payoff: an explanation of why big projects so reliably run late.

A project schedule can be drawn as a network, called an *activity network*. Its nodes are the individual activities, or tasks, that make up the project, and its links record *precedence*: a link from activity a to activity b means b cannot begin until a is finished — you cannot paint a wall before it is built. Because work only flows forward in time, the network has no loops; in the jargon it is a *directed acyclic graph*, with starting activities (no predecessors) and finishing activities (no successors). For such a network there is one object that has ruled project management since the 1950s: the *critical path*. This is the longest chain of dependent activities running from start to finish — the route through the project with no slack anywhere along it, so that delaying *any* activity on it pushes back the whole project's end date by the same amount (Kelley and Walker, 1959). The manager's time-honored prescription is therefore simple: watch the critical path, and guard it at all costs.

And yet projects keep finishing late. A survey of more than ten

thousand of them found only a few percent delivered on time and on budget, and fifteen years of improvements in management practice barely moved that figure (Flyvbjerg et al., 2003). This stubborn record hints that the critical path — a single one-dimensional thread — may be the wrong thing to watch in schedules that have grown into genuinely complex networks. Following Vazquez et al. (2023), this chapter grows such schedules from a local rule, shows that the critical path *shrinks toward irrelevance* as projects get bigger, and relocates the real source of delay risk in the network as a whole.

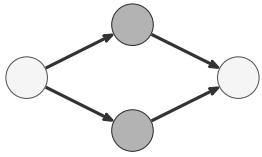
5.2 The duplication–split rule

Where does a schedule come from? Not in one piece. A plan is *refined*: a coarse activity is broken into finer ones, again and again, until the plan is detailed enough to execute. There are two natural ways to break an activity in two, and they correspond to two kinds of work. This is our local rule, with a single dial, the duplication probability q. Beginning from two activities in series, repeat: pick an existing activity at random and, with probability q, *duplicate* it, otherwise *split* it.

- *Duplication* (probability q). A *generic* activity is split into two smaller activities that run *in parallel* — think of doubling a crew so two teams do the same kind of work at once. Both copies inherit *all* the predecessors and *all* the successors of the parent. Duplication *widens* the schedule, adding parallel strands of work.

- *Split* (probability $1 - q$). A *specialized* activity is cut into two activities run *in sequence* — one specialist does the first part and hands off to another who does the second. The parent keeps its predecessors, the new activity takes over its successors, and a link joins the two. Splitting *lengthens* the schedule, extending a chain.

So the dial q measures how a project is organized (Figure 5.1). Small q means almost every refinement is a split: work is broken into ever-longer chains of specialists, and the schedule is nearly a single long line. Large q means almost every refinement is a duplication:



(a) duplication: in parallel



(b) split: in sequence

Figure 5.1: **Refining an activity.** A generic activity is replaced either by (a) two copies that run *in parallel*, both inheriting the predecessor (left, gray) and successor (right, gray) of the parent — duplication, which widens the schedule — or by (b) two specialized activities that run *in sequence* — a split, which lengthens it.

generic work is parallelized into many simultaneous strands, and the schedule fans out into a broad web. The same copying that duplicated a gene in the previous chapter here duplicates an activity; the new ingredient is the *split* — the act of *specialization* that gives the chapter its name — and the forward-flowing, start-to-finish structure of a schedule.

5.3 Predecessors and successors: identical power laws

Because the rule copies an activity together with its predecessors and successors, it is a duplication rule of the kind that produced power-law degrees in the previous chapter — only now there are two degrees to track, the number of predecessors (the in-degree) and the number of successors (the out-degree). The box works out the distribution of predecessors by following how many activities have each number of them; the same argument applies verbatim to successors. The result is in the paragraph after.

The mathematics behind it: the distribution of predecessors *(this box can be skipped)*

Let $n_k(n)$ be the number of activities with k predecessors when the schedule has n activities in all. Adding one activity changes this count by

$$n_k(n+1) = n_k(n) + q\left[\frac{k-1}{n}\,n_{k-1}(n) - \frac{k}{n}\,n_k(n) + \frac{1}{n}\,n_k(n)\right] + (1-q)\,\delta_{k,1} \tag{5.1}$$

the bracket collecting the duplication events (an activity climbs from $k-1$ to k predecessors when one of its predecessors is duplicated, and so on) and the last term adding a fresh single-predecessor activity at each split. Looking for a stable shape $n_k(n) = n\,p_k$ turns this into a recurrence,

$$p_k = q\big[(k-1)\,p_{k-1} - k\,p_k + p_k\big] + (1-q)\,\delta_{k,1}, \tag{5.2}$$

which solves in closed form (with Γ the gamma function) to

$$p_k = \prod_{s=1}^{k-1} \frac{s}{\frac{1}{q}+s}\;p_1 = \Gamma\!\left(\tfrac{1}{q}+1\right)\frac{\Gamma(k)}{\Gamma\!\left(\tfrac{1}{q}+k\right)}\,p_1, \tag{5.3}$$

with the power-law tail

$$p_k \sim k^{-1/q}. \tag{5.4}$$

The result, Eq. (5.4), is a power law with exponent $\gamma = 1/q$. A wide, duplication-heavy schedule (large q) has the heaviest tail of highly-connected activities; at $q = 1/2$ we recover the familiar value $\gamma = 2$. The very same reasoning, run on successors instead of predecessors, gives the *identical* distribution for the out-degree. That coincidence is the model's signature prediction — the spread of predecessor counts and the spread of successor counts should look the same — and it follows directly from the even-handed way duplication treats the two ends of an activity.

5.4 The vanishing critical path

The degrees are a local property; the critical path is a global one. Remarkably, the same rule fixes how the critical path grows with the

size of the project, and the answer overturns the manager's intuition. The box does the short calculation; the consequences follow.

The mathematics behind it: how the critical path scales *(this box can be skipped)*

Let c be the number of activities on the critical path of an n-activity schedule. The path can only lengthen when a *split* lands on it (a split inserts an activity in series), and the chance that the activity picked for refinement lies on the critical path is c/n. Hence

$$\frac{dc}{dn} = (1-q)\,\frac{c}{n}, \tag{5.5}$$

which integrates to $c \sim n^{1-q}$, so the *fraction* of activities on the critical path shrinks as

$$\frac{c}{n} \sim n^{-\alpha(q)}, \qquad \alpha(q) \le q. \tag{5.6}$$

Because distances here grow as a power of size, $c \sim n^{1/D}$, the network is a *fractal* with dimension

$$D = \frac{1}{1-\alpha(q)}. \tag{5.7}$$

The headline is Eq. (5.6): for any amount of duplication ($q > 0$), the critical path is a *vanishing fraction* of the project as the project grows. The bigger and more parallel the schedule, the smaller the slice of the work that the manager's cherished path actually covers — so watching the critical path means watching an ever-shrinking sliver of the whole. (Simulations confirm the power-law shrinkage, with the measured exponent α a little below the simple bound q; for instance $\alpha \approx 0.11$, 0.28, 0.43 at $q = 0.20$, 0.50, 0.70.)

The scaling also tells us what *kind* of network a schedule is. In a small-world network, distances grow only logarithmically with size — doubling the network barely lengthens the longest path. Here distances grow as a *power* of size, the mark of a *fractal* network: one whose "size" and "length" are tied together by a fixed dimension D, much as a physical fractal looks the same at every magnification. Duplication–split schedules are thus a class apart — power-law degrees like a scale-free network, but power-law (not logarithmic)

distances, unlike a small world. Specialization builds the long fractal chains; duplication stacks them up in parallel.

5.5 Real project schedules

This is not just a story about a model. Vazquez et al. (2023) examined 77 real construction projects, together worth more than ten billion dollars, which yielded 323 activity networks — snapshots of projects at different stages — ranging from about a hundred to sixteen thousand activities. Three predictions survive the encounter with data. First, the predecessor and successor counts are each power-law distributed, and fitting them *separately* yields duplication indices q that agree, just as the model demands they should. Second, those indices fall between roughly 0.1 and 0.5, most above 0.2: real schedules sit on the split-dominated, nearly-linear side — but with enough duplication to matter. Third, the share of activities on the critical path falls with project size right across the collection, confirming that the critical path really does dwindle as projects grow.

5.6 Delay risk as percolation

If the critical path fades to insignificance and yet projects are still late, where does the delay risk live? The answer is that it is spread across the whole network, and it is governed by *percolation* — the same mathematics that describes how water seeps through porous rock or how a forest fire spreads from tree to tree.

Picture an outside shock — bad weather, a strike, a pandemic — that pushes some activities past their planned finish. A delay only passes to a downstream activity if it is large enough to use up the spare time, the *free float*, built into that dependency; otherwise the buffer absorbs it and the delay stops there. Now treat each dependency as a channel that passes a delay along with some probability p, and ask what fraction of activities end up delayed beyond their own slack — the *delay incidence*. This is exactly a percolation problem: delays spread through the network of dependencies like fire through a forest, and percolation theory says such spreading has a sharp

threshold (Albert et al., 2000; Ellinas, 2019; Schwartz et al., 2002),

$$p_c = \frac{1}{\langle k \rangle}, \tag{5.8}$$

where $\langle k \rangle$ is the average number of downstream dependencies per activity. The meaning of a threshold is stark: below p_c delays stay local and project-wide overruns essentially never happen; above p_c a finite fraction of projects suffer a system-wide cascade — the fire jumps the firebreaks.

The data choose between two stories cleanly. If the critical path governed delay, the incidence ought to track $p\,c$ (transmission times critical-path size); in the real projects that correlation is negligible. If percolation governs delay, the incidence ought to switch on exactly at $p - p_c = 0$; and that is what the construction projects show — below threshold, delay incidence stays under one percent; above it, the incidence climbs, reaching fifteen percent of all activities in the worst cases. The true control knob of project delay is $p - p_c$, a property of the *whole* network, not $p\,c$, a property of one path through it.

Here is the chapter's practical lesson, and its link to the rest of the book. The very same local rule that sets the degree distributions also sets the average degree $\langle k \rangle$, and through Eq. (5.8) the threshold for catastrophic delay. A project's tendency to run late is written into the network its activities have grown into — not into any single path through it.

5.7 Summary

Project schedules are forward-flowing networks of activities, grown by a local refinement rule: pick an activity and either duplicate it into parallel copies (probability q) or split it into specialized sequential parts. This duplication–split rule predicts identical power-law distributions for the numbers of predecessors and successors, $p_k \sim k^{-1/q}$; a critical path that grows only as $c \sim n^{1-q}$ and so vanishes as a fraction of large projects; and a fractal, rather than small-world, geometry. All three hold in 77 real construction projects. And because the critical path becomes irrelevant at scale, delay risk is

better read through percolation, with a threshold $p_c = 1/\langle k \rangle$ fixed by the grown network's average degree — so that project performance, in the end, is a property of the activity network as a whole.

Chapter 6

Emergence of communities

6.1 Communities nobody designed

So far the local rules of this book have explained the *degree* structure of networks — their power laws, their clustering hierarchy, their correlations. This chapter turns to a feature of an altogether different kind: *community structure*. A community is a group of nodes that are densely connected among themselves but only sparsely connected to the rest of the network — a clump that hangs together. They are everywhere in real data: people fall into social circles, web pages into topics, proteins into functional modules, each group talking mostly to itself. Splitting a network into such groups is what researchers mean by *community detection*, and a network that splits cleanly is said to have strong community structure.

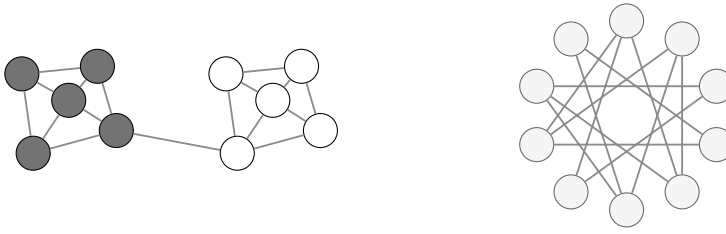
Where do these communities come from? The usual answer points to differences between the nodes themselves — *node heterogeneity*, in the jargon. People of the same profession or politics, web pages on the same subject, proteins with the same biological function: nodes of a kind, the story goes, are simply more likely to link to one another, and that built-in preference is what carves the network into groups. On this view, communities are stamped onto the network by properties the nodes carried in from the start.

The thesis of this chapter, following Vazquez (2025), is that no such stamp is needed. *Communities emerge from local growth rules even when every node is identical.* The first hint is decades old: Jin et al. (2001) showed that a social network grown purely by triadic closure — the friends-of-friends rule of Chapter 3 — develops communities although its nodes are all of one type, and the same was later seen in the gene-duplication models of Chapter 4 (Solé et al., 2002). If a featureless local rule segregates a network into groups all by itself, then community structure joins preferential attachment and the clustering hierarchy as one more thing that *emerges* from locality rather than being designed in.

To make that claim sharp we should ask it as a quantitative question, and the right way to frame it comes from a corner of mathematics called *Ramsey theory.* Ramsey theory studies how large a random structure must become before some orderly pattern is forced to appear. Its best-known nugget is a party fact: in any gathering of six people, there are always either three who all know one another or three who are all mutual strangers — order of that kind is unavoidable once the group is big enough. We will ask the analogous question of growing networks: given a local growth rule, how large must the network become before it is virtually *certain* to contain communities? This chapter is, in effect, a Ramsey theory of community structure in growing networks.

6.2 The Ramsey community number

Before we can measure when communities appear, we need a trustworthy way to *detect* them, and here some care is required, because the most popular method can be fooled. That method, *modularity optimization*, looks for the division of a network that packs the most links inside groups; the trouble is that it will happily report groups even in a completely random network that has no real structure at all (Newman, 2006; Fortunato, 2010). To avoid being deceived, the work relies instead on two more discerning tools: the *stochastic block model*, which fits the network to a statistical model of blocks and keeps only the structure that genuinely earns its keep (Karrer and Newman, 2011), cross-checked against *Infomap*, which finds groups by watching where a random walker gets trapped (Rosvall



(a) grown by a local rule (b) after random rewiring

Figure 6.1: **Communities emerge from local rules.** (a) A network grown by a local rule settles into dense communities (two here, one dark gray, one white) joined by only a few links. (b) Shuffling the links at random while keeping every node's number of connections destroys the communities — so the structure was the work of the rule, not of the degrees alone.

and Bergstrom, 2008). Both, crucially, report *no* communities in a truly random network.

That last point also supplies the essential control. To be sure a community is the work of the growth rule and not merely of the degrees, one compares the grown network against a *rewired* version of itself — the links shuffled at random while every node keeps its number of connections (this shuffling is called the configuration model). If the communities survive the shuffle, they were never more than an accident of the degrees; if they vanish, they were the doing of the local rule. For every local model in this chapter, the shuffled version has no communities at all — so whatever structure we find is genuinely built by the dynamics.

With detection settled, we can pose the question precisely (Figure 6.1). The box below states the two definitions that do the work; both are short, and their meaning is given in words right after, so the box can be skipped.

The mathematics behind it: the communities likelihood and the Ramsey community number *(this box can be skipped)*

For a growth rule f_G that produces a network of n nodes, let a detection method f_κ return the number κ of communities it finds. Define the *communities likelihood* as the probability — estimated by generating many networks — that an n-node instance has more than one community,

$$P_\kappa(n) = \text{Prob}\{\, \kappa \geq 2 \,\}. \tag{6.1}$$

The *Ramsey community number* is then the smallest size at which communities are all but guaranteed, at a confidence $1 - \epsilon$ (we use $\epsilon = 0.05$, i.e. 95%):

$$r_\kappa = \min\{\, n : P_\kappa(n) \geq 1 - \epsilon \,\}. \tag{6.2}$$

A rule is said to have the *emergent communities property* when such an r_κ exists.

In plain terms: $P_\kappa(n)$ is the chance that a network grown to size n has communities, and the Ramsey community number r_κ is the size beyond which that chance is essentially $(1 - \epsilon)100\%$. A *finite* r_κ is the fingerprint of emergence — grow the network past it and communities are a near-certainty. A rule for which no such size exists, no matter how large the network, simply does not make communities on its own. The whole question "do communities emerge from this rule?" thus reduces to "is r_κ finite?"

6.3 Local rules grow communities

We now line up two local rules against two controls. The local rules — local search and duplication–split, both met earlier in this book — have the emergent communities property; the controls — reshuffling a network's links, and growing one with no locality at all — do not. Holding the four side by side is the whole argument of the chapter in miniature.

Local search. The model LS(n, d) is the search rule of Chapter 2 in undirected form (Vazquez, 2001): each new node takes a short walk of d steps from a random starting node and links to every node

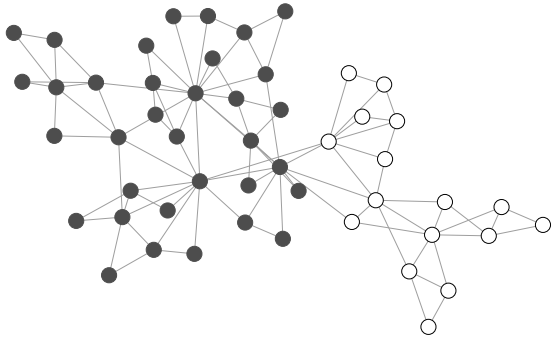


Figure 6.2: **Communities in a local-search network.** An instance of the local-search model LS(n=50, d=1). The stochastic block model resolves two communities, shown dark gray and white; they emerge even though the rule treats every node alike and never mentions groups.

it visits. It carries preferential attachment and high clustering, and — though nothing in the rule ever mentions groups — the block model resolves clear communities, with P_κ climbing from near zero around $n = 10$ to a near-certain one beyond about $n = 100$, and staying there out to a hundred thousand nodes (Vazquez, 2025). Figure 6.2 shows one such network.

Duplication–split. The model DS(n, q) is the undirected cousin of the project-schedule rule of Chapter 5: a new node either copies all of a random node's links (probability q) or splits one of its links in two. This model is the decisive case, because it builds *no triangles whatsoever* — duplication makes none, and splitting destroys any that might form — and yet communities emerge just as reliably. The lesson is sharp: a network does not need to be cliquish to have communities. What it needs is locality, not triangles. Figure 6.3 shows an instance: two communities, and not a triangle in sight.

Degree-preserving randomization. To be sure the communities are the work of the local rule and not of the degrees alone, we need a way to wipe out the local structure while leaving every degree

untouched. The *configuration model*, met when we defined the Ramsey number, does exactly that: it takes a finished network and reshuffles its links at random, keeping each node's number of connections fixed but scrambling *which* nodes are joined. Applied to a local-search or a duplication–split network, this reshuffling destroys the communities outright — the randomized versions have no finite Ramsey number at all (the two "randomized" rows in the table below). Whatever the local rules built, the bare list of degrees cannot sustain.

Barabási–Albert. Finally, BA(n, m) — the canonical preferential-attachment model — is our *non-local* control: a growing, rich-get-richer network in which each newcomer attaches its m links to nodes drawn from the *whole* graph, never once consulting a neighborhood. Despite its growth and its rich-get-richer rule it has no emergent communities (the last row of the table), and the next section is devoted to teasing apart exactly why.

A few representative Ramsey numbers, based on numerical estimates for $\epsilon = 0.05$ (Vazquez, 2025), make the pattern concrete:

Model	r_κ (block model)	r_κ (Infomap)
LS$(d = 1)$	81	222
DS$(q = 0.5)$	98	622
LS$(d = 1)$, randomized	—	—
DS$(q = 0.5)$, randomized	—	—
BA$(m = 2)$	—	—

Every genuinely local rule has a finite Ramsey number — grow it past a few hundred nodes and communities are certain. The rules without local structure — the bottom three rows, the two reshuffled networks and the Barabási–Albert model — have none: the dashes mean communities never become certain, however large the network. The exact figures differ between the two detection methods — there is no single "true" definition of a community, and where the line falls depends a little on how one draws it — but the *split* between local rules (finite) and non-local ones (infinite) is the same either way. Local structure is a necessary condition for communities to emerge.

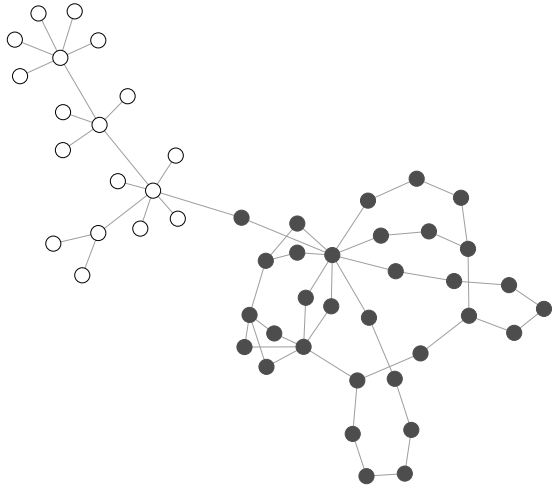


Figure 6.3: **Communities in a duplication–split network.** An instance of the duplication–split model DS(n=50, q=0.3). Two communities (dark gray and white) again emerge — even though this rule builds no triangles at all, so the network is far less cliquish than the local-search one of Figure 6.2.

6.4 When globality fails: the Barabási–Albert model

The natural control is the Barabási–Albert model (Barabási and Albert, 1999), the textbook growing network with *no* locality at all: each newcomer attaches its m links to nodes chosen from the *whole* network in proportion to their degree. It has growth and preferential attachment, but no step that ever looks at a node's neighborhood — and it does *not* have the emergent communities property. Its community likelihood rises only to about one-half near a thousand nodes and then *falls* as the network grows further, never reaching certainty. The little structure it has comes from the mere age-ordering of its nodes, not from any local rule, and it washes out at scale.

The boundary case is illuminating. Give the model a baseline

"attractiveness" a, so that a node is chosen with a probability proportional to $a + k_i / \sum_j k_j$ — a fixed head-start a on top of the usual degree preference. This tunes the degree exponent to $\gamma = 2 + a/m$. Tracking the Ramsey number across these variants reveals a clean rule of thumb: communities emerge exactly when the network is *clustered* enough. The Ramsey number falls as the average clustering coefficient rises, and *diverges* — communities never become certain — as the clustering goes to zero (Vazquez, 2025). The standard Barabási–Albert model sits right at that vanishing point, because its clustering decays toward zero as it grows (like $(\ln n)^2/n$): the bigger it gets, the fewer short cycles it has, and short cycles are what local structure is made of. So even here the principle holds — where the dynamics happen to generate appreciable clustering (an *implicit* bit of local structure), communities appear; where clustering fades, they do not.

6.5 How local must "local" be?

If locality is the cause, exactly how local must a connection be to count? Two dials settle it: the length of the cycles a rule creates, and the distance it reaches across. Triadic closure closes triangles (cycles of length three); the duplication–split rule works with squares (length four). To sweep the cycle length freely, use the *bubble model*: at each step a chain of L new nodes is added and its two ends are attached to the two ends of a randomly chosen existing link, creating a fresh cycle of length $L + 2$. As L grows — as the rule weaves ever longer cycles — the Ramsey number climbs with it, and steeply: from $r_\kappa = 108$ for the shortest cycles to more than fifty thousand by $L = 9$ (Vazquez, 2025), as the table below records. The climb is not quite smooth — riding on top of it is a regular even–odd alternation, the odd-length cycles bringing communities on sooner than the even cycles around them, an effect plainest for the shorter cycles — but the overall message is unmistakable: short cycles, tight local feedback, bring communities on soonest.

L	cycle length $L+2$	r_κ
1	3	108
2	4	326
3	5	246
4	6	2201
5	7	752
6	8	6595
7	9	2740
8	10	15863
9	11	56160

Distance matters just as much. In the bubble model the chain's two ends attach to nodes that are a single step apart (the ends of one link). Loosen this so the ends may attach to nodes up to W steps apart, with the chain length fixed at $L = 1$. For $W = 1$ communities still emerge with certainty, at $r_\kappa = 108$; stretching to $W = 2$ pushes the threshold up roughly tenfold, to $r_\kappa = 1090$, but the communities still come. At $W = 3$ — and likewise if the ends are joined to nodes picked anywhere at random — they stop appearing altogether, the Ramsey number ceasing to exist (the dash in the table) (Vazquez, 2025).

W	r_κ
1	108
2	1090
3	—

For this model, "local" means roughly two steps: communities form when new links reinforce ties among nearby nodes, and dissolve once links are allowed to leap across the network. This is the whole book's premise made quantitative. It is not growth, nor preferential attachment, nor even clustering in the abstract that carves a network into communities — it is the *locality* of the rule by which it grew.

6.6 Summary

Community structure, long blamed on built-in differences between nodes, is instead something that *emerges* from growing a network

by local rules, even when every node is identical. Made precise through the Ramsey community number r_κ — the size beyond which communities are all but certain — the statement is clean: the local rules of earlier chapters (search and duplication–split) have a finite r_κ, while their degree-preserving shuffles, and the non-local Barabási–Albert model, have none. Triangles are not required (duplication–split makes none and still segregates), but locality is: communities come soonest for rules that close short cycles between nearby nodes, and disappear once links may span the whole network. Having already accounted for degree distributions, the clustering hierarchy, and degree correlations, this chapter adds community structure to the list of global patterns that locality alone can explain (Vazquez, 2002).

Chapter 7

Boosting shortest path multiplicity

7.1 The redundancy of shortest paths

The small-world property, met back in the introduction, promises that any two nodes of a real network are joined by a *short* path — only a few hops apart. It says nothing, though, about how *many* short paths there are. That count turns out to matter a great deal. If the only shortest route between two nodes runs through a single chain of intermediaries, the connection is fragile: cut one link and the shortest route is gone. If instead there are several shortest routes side by side, the connection is robust, and whatever flows between the two nodes — a message, a signal, a delay — can travel along all of them at once. The quantity that captures this is the *shortest path multiplicity*: the number of distinct shortest routes between a pair of nodes. A simple picture: if two friends can each be reached from you in two handshakes, but by two completely different chains of acquaintances, the multiplicity of that connection is two. In real networks this multiplicity is conspicuously high.

And it is not an isolated fact. Deng et al. (2026) found that, across many real networks, the abundance of multiple shortest paths is strongly correlated with *community structure* — the very groupings of the previous chapter — more strongly than with any other

standard measure. Two networks of the same size and the same degree distribution can differ widely in their path multiplicity, and the difference tracks how community-rich they are. A correlation like that begs for an explanation. Does community structure somehow *cause* paths to multiply, or do both spring from a common deeper source? The previous chapter hands us the obvious suspect: local growth rules, which we saw conjure communities out of nothing but locality. This closing chapter, following Vazquez (2026), completes the case by showing that those same local rules also *boost* the multiplicity of shortest paths — so that the striking link between paths and communities is not one causing the other, but the shared shadow cast by the rule that grew the network.

The mechanism is easy to picture once named. Local rules build short *cycles* — closed loops — as a matter of course: a triangle whenever a friend introduces a friend, a square whenever a gene is copied. And a cycle is exactly a structure that offers more than one way around. In a network with no loops at all — a *tree* — there is precisely one route between any two nodes; the more loops a network carries, the more alternative shortest routes it affords. Path multiplicity is simply what local rules leave behind when they close loops.

7.2 Defining the multiplicity

We need only one quantity. Write $\mu_{ij}(G)$ for the number of distinct shortest paths between nodes i and j in a network G — counting every route that ties for the minimum number of hops. Averaging this over all pairs of nodes, and over many networks grown the same way, gives the *average shortest path multiplicity* $\langle \mu \rangle$, the single number we will follow as the network grows to size n. A tree, with its unique routes, has $\langle \mu \rangle = 1$ exactly; anything above one measures the redundancy that loops provide. The whole chapter is about one question: how fast does $\langle \mu \rangle$ rise as the network gets bigger, and does that rate betray whether the network was grown by a local rule?

7.3 The baseline: random networks

To judge what counts as "boosted," we first need the expectation in the *absence* of any local structure. The natural reference is a random network. For a dense random network — the classic Erdős–Rényi graph, in which every possible link is present with some fixed probability — the average multiplicity grows in proportion to the size n (Dong et al., 2025). But real networks are not dense; they are *sparse*, each node having only a handful of links no matter how large the network. The fair baseline is therefore a *sparse* random network, which we make by taking a network grown by a local rule and then shuffling its links at random while leaving each node's number of links untouched — the same configuration-model scramble used as a control in the previous chapter. Stripped of their local order this way, these sparse random networks have a multiplicity that creeps up only *logarithmically* with size, as determined by numerical simulations (Vazquez, 2026):

$$\langle \mu \rangle_{S,0} = a + b \log n, \tag{7.1}$$

the subscript $S, 0$ marking a network that is sparse (S) and unstructured (0). Logarithmic growth is about as slow as growth gets — to add a fixed amount of multiplicity you must *multiply* the network's size, not just add to it. That sluggish, degrees-only baseline is the bar that local rules must clear. Anything that grows faster is the doing of the rule.

Of note, the Barabási–Albert model falls on this same baseline. Although it is a *growing* network rather than a randomized one, it has no local structure — newcomers attach across the whole graph, never consulting a neighborhood — and its shortest path multiplicity creeps up only logarithmically, exactly as the reshuffled networks do. This echoes the previous chapter, where the very same model failed to develop communities: lacking locality, it builds neither the modules nor the redundant paths, and on both counts it behaves as though its links had been shuffled at random.

7.4 Local rules boost the multiplicity

Now run the very same local models of the previous chapter *without* shuffling — the local-search rule, the duplication–split rule, the bubble model — and the multiplicity climbs decisively faster than the logarithmic baseline. It helps to keep a ladder of growth rates in mind: logarithmic (slowest, the baseline), then *log-quadratic* (faster), then *exponential* (explosive, fastest of all). The local rules fall into two of these rungs, as deduced from numerical simulations (Vazquez, 2026).

Most of them — local search and the bubble model among them — show *log-quadratic* growth,

$$\langle \mu \rangle_{S,I} = a + b \log n + c(\log n)^2, \tag{7.2}$$

the extra squared term lifting the curve clearly above the straight logarithmic baseline. Call this type I.

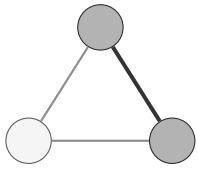
The duplication–split rule is different, and dramatically so: its multiplicity grows *exponentially* with size,

$$\langle \mu \rangle_{S,II} = a\, e^{bn}, \qquad a, b > 0, \tag{7.3}$$

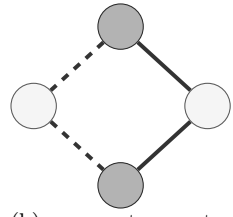
which we call type II. Exponential growth is a different world from logarithmic — the redundancy proliferates beyond anything the other rules manage. Copying, it turns out, is the great multiplier of shortest paths, and that is what the title of this chapter names.

7.5 Cycles and the boost

Why should duplication outrun the rest? The answer is in the kind of loop each rule makes. The triadic-closure rule of the friends chapter closes *triangles*, loops of length three — but in a triangle the two nodes flanking a third are themselves directly linked, so there is still only *one* shortest route between any pair. A triangle adds cohesion but no extra routes. Duplication, by contrast, closes *squares*, loops of length four, and a square is the smallest loop that genuinely offers a choice: its two opposite corners are two hops apart by *two* different paths, one around each side. That single observation is the seed



(a) triangle: one route



(b) square: two routes

Figure 7.1: **Why squares multiply routes.** The two highlighted corners are the endpoints. (a) In a triangle they are joined directly — a single shortest route. (b) In a square they are opposite corners, two steps apart by *two* distinct routes (drawn solid and dashed), one around each side. Local rules that close four-cycles, such as gene duplication, thread the network with these redundant shortest paths.

of the boost; the box pushes it a little further, and the takeaway follows.

Figure 7.1 makes the contrast concrete.

The mathematics behind it: why copying multiplies routes *(this box can be skipped)*

When a node is duplicated, the copy and the original share all their neighbors, so any two of those neighbors sit at opposite corners of a four-cycle — each such pair instantly gains a second shortest route. A duplicated node of degree k creates $\binom{k}{2} = k(k-1)/2$ such squares in a single stroke. Moreover the duplication rule carries its own preferential attachment, the chance of copying a node being $\pi_k = qk/n$, which (as in Chapter 4) builds a heavy-tailed degree distribution with exponent $\gamma = 1 + 1/q$. Squares are thus laid down in proportion to $k(k-1)$ on top of a distribution with many high-degree nodes — a compounding that is the suspected origin of the exponential law, Eq. (7.3), though a full account of it remains open.

The lesson of the box is that squares, not triangles, are what multiply shortest routes: each duplication endows many pairs of nodes with a second equally-short path, and because high-degree

nodes (which spawn the most squares) are themselves abundant, the redundancy compounds. The bubble model, which lets us dial the loop length directly — a chain of L added nodes makes a loop of length $L + 2$ — confirms the geometric intuition: even-length loops, whose opposite corners enjoy two equal routes, enrich the network more than odd ones. (It is a useful caution, though, that the bubble model's multiplicity still grows only log-quadratically, type I, whatever the loop length; the square-counting explains why duplication *beats* triadic closure, but the full exponential ferocity of duplication–split needs the compounding above.)

One more clue shows that randomness itself plays a role. The Dorogovtsev–Goltsev–Mendes "pseudofractal" network (Dorogovtsev et al., 2002) is a *deterministic* relative of the bubble model — grown by a fixed, clockwork rule rather than a random one. Its multiplicity does *not* accelerate into the log-quadratic regime; it grows only log-linearly, like the shuffled baseline. So the speed-up in Eq. (7.2) is not a property of the loop geometry alone — it requires the *stochastic* character of the growth. It is the interplay of randomness and locality together, not either by itself, that boosts the multiplicity.

7.6 Two sides of one coin

Because local rules also create communities, and the number of communities grows along with the network, we can ask how the multiplicity depends on the number of communities directly. The relationship is a clean one — linear for every local model tested, and (once again the exception) exponential for duplication–split. This is the empirical correlation of Deng et al. (2026) restated from the inside: more communities go hand in hand with more shortest paths.

But the point of the chapter is that this partnership is not a case of one causing the other. "Local rules boost shortest path multiplicity" and "communities accompany high shortest path multiplicity" are two true statements about the same networks, and as descriptions they say the same thing. Only one of them, however, names a *mechanism*. Communities are something an algorithm reads off a finished network after the fact; the local rule is the process that actually built that network, node by node, closing the very loops that both carve it into communities and thread it with redundant

routes. Community structure and path multiplicity are two sides of a single coin — and the coin is locality.

7.7 Summary

The shortest paths of real networks are richly redundant, and that redundancy travels with their community structure. Sparse random networks, having no local order, accumulate multiplicity only logarithmically as they grow; networks built by local rules accumulate it far faster — log-quadratically for most rules, and exponentially for duplication, which closes the four-cycles that give opposite corners two shortest routes apiece. Since the same local rules also produce the communities, the observed tie between path multiplicity and communities is best understood not as cause and effect but as two consequences of one underlying local dynamics.

Chapter 8

Outlook

8.1 One principle, many faces

We have now followed a single idea through six concrete settings, and it is worth gathering the threads before asking what lies beyond them. In every case the network grew by a *local* rule — a newcomer acting only on a node it had reached and that node's immediate neighbors — and in every case the same family of global regularities emerged unbidden.

Searchable networks, such as the citation graph and the World Wide Web, grow by exploration: a newcomer surfs along links and connects to what it finds. Modeling the surfer as the very random walk that defines Google's PageRank, we found that a page's visiting probability rises in a straight line with its in-degree, so that the local act of walking manufactures an effective "rich get richer." From it follow a power-law degree distribution whose exponent is fixed purely by how the network is explored, a clustering hierarchy in which cliquishness falls with degree, and generically disassortative correlations (Chapter 2).

Social networks grow by introduction: a friend introduces a friend, closing a triangle. Formalizing the friend-of-a-friend as a *potential link*, born in proportion to degree, we recovered the same rich-get-richer and the same clustering hierarchy — but now with *assortative* correlations, the well-connected keeping one another's

company (Chapter 3). The contrast with the searchable networks is that chapter's quiet lesson: the sign of the degree correlations is not a free parameter but a readout of whether the rule reaches outward, to the periphery, or inward, within a neighborhood.

Biological networks grow by copying: a gene is duplicated, its protein inheriting all of the ancestor's interactions, and the copies then diverge. Because a node gains a link whenever one of its neighbors is duplicated, copying too is a rich-get-richer engine; self-interactions seed the clustering hierarchy, and coupling divergence to duplication adds a density transition and a multifractal degree distribution born of inheritance, alongside the disassortative mixing measured in real protein networks (Chapter 4). Here the three mechanisms met in one sentence: a rule that attaches to a *neighbor* of a randomly chosen node attaches in proportion to degree, because a well-connected node simply belongs to more neighborhoods. Preferential attachment, the clustering hierarchy, and the degree correlations are all the generic harvest of locality.

The same copying, carried into project schedules, refines each activity either by duplicating it into parallel copies or by splitting it into specialized sequential parts. The rule predicts identical power laws for an activity's predecessors and successors, a critical path that vanishes as a fraction of large projects, and a fractal rather than small-world geometry — all borne out in seventy-seven real construction projects — and it relocates the risk of delay from the celebrated critical path to a percolation threshold set by the whole network's average degree (Chapter 5).

Finally, locality proved to shape structure above the level of the individual node. The same rules that set the degrees also segregate the network into communities — even when every node is identical — a fact we made precise through the Ramsey community number, the size beyond which communities are all but certain (Chapter 6); and they thread the network with redundant shortest paths, multiplying the routes between nodes far faster than chance and in lockstep with those very communities (Chapter 7).

And with that, the book comes full circle. We set out asking where the global regularities of real networks come from — their power-law degrees, their small-world distances, their clustering hierarchy, their degree correlations — and answered, one mechanism

at a time, that none of them need be imposed from above: each emerges from a way of growing in which no node ever uses more than local knowledge. A node that surfs to a neighbor, befriends a friend, is copied from an ancestor, or refines a task asks nothing of the network as a whole — and yet the networks these humble acts build carry, in the aggregate, the full and unmistakable signature of real complex systems.

8.2 What remains to be done

If the thesis of this book is right, it opens more questions than it closes.

Several are concrete. The exponential growth of shortest path multiplicity under the duplication–split rule — so much faster than the log-quadratic growth of every other rule we studied — still wants a full explanation; the square-counting argument shows why copying outruns triadic closure, but not why it runs away exponentially. The Ramsey community number, likewise, remains tied to the community-detection method used to measure it, and a notion of emergence that does not depend on the algorithm would put the whole idea on firmer ground. Even the self-organized critical state of the exhaustive searcher, where a scale-free network appears across a whole range of the control parameter rather than at one fine-tuned value, deserves to be connected more tightly to the wider physics of criticality.

Others are programmatic. Throughout, we read each local rule *forward*, from mechanism to statistics; the empirical chapters hint at the reverse problem — reading the rule *backward*, inferring from a single observed network which local moves grew it, and with what parameters, as we did when fitting a duplication index to project schedules. A general inverse theory, turning measured correlations, clustering hierarchies and community structure into an estimate of the underlying rule, would make these models genuinely diagnostic. And real networks are surely not grown by one rule in isolation: a social graph is part introduction and part search, a cellular network part duplication and part fresh interaction. How competing local moves combine — whether their signatures simply add, or interfere to produce something neither would alone — is largely open.

The setting, too, can be widened. The networks of this book are

mostly unweighted, undirected and static, whereas real ones carry weights, directions, timestamps, and several kinds of link at once. Whether the locality principle survives into weighted, temporal and multilayer networks — and what new regularities it might generate there — is a natural frontier. So is the coupling of growth to *function*: the percolation of delays through a project schedule was a first glimpse of a dynamics running on a network that local rules had built, and the interplay between how a network grows and what it then does promises more than either studied alone.

None of these questions threatens the central claim; each extends it. Local rules are the driving mechanism; everything else follows.

Bibliography

Réka Albert and Albert-László Barabási. Statistical mechanics of complex networks. *Reviews of Modern Physics*, 74(1):47–97, 2002. doi: 10.1103/RevModPhys.74.47.

Réka Albert, Hawoong Jeong, and Albert-László Barabási. Error and attack tolerance of complex networks. *Nature*, 406(6794): 378–382, 2000. doi: 10.1038/35019019.

Albert-László Barabási and Réka Albert. Emergence of scaling in random networks. *Science*, 286(5439):509–512, 1999. doi: 10.1126/science.286.5439.509.

Sergey Brin and Lawrence Page. The anatomy of a large-scale hypertextual Web search engine. *Computer Networks and ISDN Systems*, 30(1–7):107–117, 1998. doi: 10.1016/S0169-7552(98) 00110-X.

Jörn Davidsen, Holger Ebel, and Stefan Bornholdt. Emergence of a small world from local interactions: Modeling acquaintance networks. *Physical Review Letters*, 88(12):128701, 2002. doi: 10.1103/PhysRevLett.88.128701.

Ye Deng, Jun Wu, Xin Lu, Petter Holme, Daqing Li, Zengru Di, Guanrong Chen, and Jürgen Kurths. Community structure unveils the path multiplicity in complex networks. *Nature Communications*, 17:2283, 2026. doi: 10.1038/s41467-026-70369-4.

Yu Dong, Ye Deng, and Jun Wu. Bimodal distribution of path multiplicity in random networks. *Chaos, Solitons & Fractals*, 193: 116124, 2025. doi: 10.1016/j.chaos.2025.116124.

S. N. Dorogovtsev, A. V. Goltsev, and J. F. F. Mendes. Pseudofractal scale-free web. *Physical Review E*, 65(6):066122, 2002. doi: 10.1103/PhysRevE.65.066122.

Sergey N. Dorogovtsev and José F. F. Mendes. Evolution of networks. *Advances in Physics*, 51(4):1079–1187, 2002. doi: 10.1080/00018730110112519.

Christos Ellinas. The domino effect: An empirical exposition of systemic risk across project networks. *Production and Operations Management*, 28(1):63–81, 2019. doi: 10.1111/poms.12890.

Michalis Faloutsos, Petros Faloutsos, and Christos Faloutsos. On power-law relationships of the Internet topology. *ACM SIGCOMM Computer Communication Review*, 29(4):251–262, 1999. doi: 10.1145/316194.316229.

Bent Flyvbjerg, Mette K. Skamris Holm, and Søren L. Buhl. How common and how large are cost overruns in transport infrastructure projects? *Transport Reviews*, 23(1):71–88, 2003. doi: 10.1080/01441640309904.

Allan Force, Michael Lynch, F. Bryan Pickett, Angel Amores, Yilin Yan, and John Postlethwait. Preservation of duplicate genes by complementary, degenerative mutations. *Genetics*, 151(4):1531–1545, 1999. doi: 10.1093/genetics/151.4.1531.

Santo Fortunato. Community detection in graphs. *Physics Reports*, 486(3–5):75–174, 2010. doi: 10.1016/j.physrep.2009.11.002.

Michelle Girvan and M. E. J. Newman. Community structure in social and biological networks. *Proceedings of the National Academy of Sciences*, 99(12):7821–7826, 2002. doi: 10.1073/pnas.122653799.

Mark S. Granovetter. The strength of weak ties. *American Journal of Sociology*, 78(6):1360–1380, 1973. doi: 10.1086/225469.

Petter Holme and Beom Jun Kim. Growing scale-free networks with tunable clustering. *Physical Review E*, 65(2):026107, 2002. doi: 10.1103/PhysRevE.65.026107.

Hawoong Jeong, Sean P. Mason, Albert-László Barabási, and Zoltán N. Oltvai. Lethality and centrality in protein networks. *Nature*, 411(6833):41–42, 2001. doi: 10.1038/35075138.

Emily M. Jin, Michelle Girvan, and M. E. J. Newman. Structure of growing social networks. *Physical Review E*, 64(4):046132, 2001. doi: 10.1103/PhysRevE.64.046132.

Brian Karrer and M. E. J. Newman. Stochastic blockmodels and community structure in networks. *Physical Review E*, 83(1):016107, 2011. doi: 10.1103/PhysRevE.83.016107.

James E. Kelley and Morgan R. Walker. Critical-path planning and scheduling. In *Papers presented at the December 1–3, 1959, Eastern Joint IRE-AIEE-ACM Computer Conference*, pages 160–173, 1959. doi: 10.1145/1460299.1460318.

J. Kim, P. L. Krapivsky, B. Kahng, and S. Redner. Infinite-order percolation and giant fluctuations in a protein interaction network. *Physical Review E*, 66(5):055101, 2002. doi: 10.1103/PhysRevE.66.055101.

P. L. Krapivsky and S. Redner. Organization of growing random networks. *Physical Review E*, 63(6):066123, 2001. doi: 10.1103/PhysRevE.63.066123.

Sergei Maslov and Kim Sneppen. Specificity and stability in topology of protein networks. *Science*, 296(5569):910–913, 2002. doi: 10.1126/science.1065103.

M. E. J. Newman. The structure of scientific collaboration networks. *Proceedings of the National Academy of Sciences*, 98(2):404–409, 2001. doi: 10.1073/pnas.98.2.404.

M. E. J. Newman. Assortative mixing in networks. *Physical Review Letters*, 89(20):208701, 2002. doi: 10.1103/PhysRevLett.89.208701.

M. E. J. Newman. The structure and function of complex networks. *SIAM Review*, 45(2):167–256, 2003. doi: 10.1137/S003614450342480.

M. E. J. Newman. Modularity and community structure in networks. *Proceedings of the National Academy of Sciences*, 103(23):8577–8582, 2006. doi: 10.1073/pnas.0601602103.

Susumu Ohno. *Evolution by Gene Duplication*. Springer-Verlag, Berlin, 1970.

Gergely Palla, Imre Derényi, Illés Farkas, and Tamás Vicsek. Uncovering the overlapping community structure of complex networks in nature and society. *Nature*, 435(7043):814–818, 2005. doi: 10.1038/nature03607.

Romualdo Pastor-Satorras, Alexei Vazquez, and Alessandro Vespignani. Dynamical and correlation properties of the Internet. *Physical Review Letters*, 87(25):258701, 2001. doi: 10.1103/PhysRevLett.87.258701.

Derek J. de Solla Price. Networks of scientific papers. *Science*, 149 (3683):510–515, 1965. doi: 10.1126/science.149.3683.510.

Erzsébet Ravasz and Albert-László Barabási. Hierarchical organization in complex networks. *Physical Review E*, 67(2):026112, 2003. doi: 10.1103/PhysRevE.67.026112.

Erzsébet Ravasz, Anna Lisa Somera, Dale A. Mongru, Zoltán N. Oltvai, and Albert-László Barabási. Hierarchical organization of modularity in metabolic networks. *Science*, 297(5586):1551–1555, 2002. doi: 10.1126/science.1073374.

Martin Rosvall and Carl T. Bergstrom. Maps of random walks on complex networks reveal community structure. *Proceedings of the National Academy of Sciences*, 105(4):1118–1123, 2008. doi: 10.1073/pnas.0706851105.

N. Schwartz, R. Cohen, D. ben Avraham, A.-L. Barabási, and S. Havlin. Percolation in directed scale-free networks. *Physical Review E*, 66(1):015104, 2002. doi: 10.1103/PhysRevE.66.015104.

Sara Nadiv Soffer and Alexei Vazquez. Network clustering coefficient without degree-correlation biases. *Physical Review E*, 71(5):057101, 2005. doi: 10.1103/PhysRevE.71.057101.

Ricard V. Solé, Romualdo Pastor-Satorras, Eric Smith, and Thomas B. Kepler. A model of large-scale proteome evolution. *Advances in Complex Systems*, 5(1):43–54, 2002. doi: 10.1142/S021952590200047X.

Alexei Vazquez. Knowing a network by walking on it: emergence of scaling. Preprint, arXiv:cond-mat/0006132, 2000. Available at `https://arxiv.org/abs/cond-mat/0006132`.

Alexei Vazquez. Disordered networks generated by recursive searches. *Europhysics Letters*, 54(4):430–435, 2001. doi: 10.1209/epl/i2001-00259-y.

Alexei Vazquez. *Degree correlations and clustering hierarchy in networks: measures, origin and consequences.* Phd thesis, SISSA, 2002. Available at `https://www.researchgate.net/publication/395472157`.

Alexei Vazquez. Growing network with local rules: Preferential attachment, clustering hierarchy, and degree correlations. *Physical Review E*, 67(5):056104, 2003. doi: 10.1103/PhysRevE.67.056104.

Alexei Vazquez. Emergence of network communities driven by local rules. *Physical Review E*, 111(6):064314, 2025. doi: 10.1103/PhysRevE.111.064314.

Alexei Vazquez. Local network evolution rules drive shortest path multiplicity. Preprint, arXiv:2605.25237, 2026. Available at `https://arxiv.org/abs/2605.25237`.

Alexei Vazquez, Romualdo Pastor-Satorras, and Alessandro Vespignani. Large-scale topological and dynamical properties of the Internet. *Physical Review E*, 65(6):066130, 2002. doi: 10.1103/PhysRevE.65.066130.

Alexei Vazquez, Alessandro Flammini, Amos Maritan, and Alessandro Vespignani. Modeling of protein interaction networks. *ComPlexUs*, 1(1):38–44, 2003. doi: 10.1159/000067642.

Alexei Vazquez, Iacopo Pozzana, Georgios Kalogridis, and Christos Ellinas. Activity networks determine project performance. *Scientific Reports*, 13:509, 2023. doi: 10.1038/s41598-022-27180-0.

Andreas Wagner. The yeast protein interaction network evolves rapidly and contains few redundant duplicate genes. *Molecular Biology and Evolution*, 18(7):1283–1292, 2001. doi: 10.1093/oxfordjournals.molbev.a003913.

Duncan J. Watts and Steven H. Strogatz. Collective dynamics of 'small-world' networks. *Nature*, 393(6684):440–442, 1998. doi: 10.1038/30918.